

**Treatment Success and Failure Prediction Using Machine Learning Models  
for Pulmonary Tuberculosis Patients Under the Department of Health-Cordillera  
Administrative Region (DOH-CAR) Tuberculosis-Directly Observed Treatment,  
Short-course (TB-DOTS) Program**

Aaron Kyle M. Aguilar, Mark Lestat C. Agustin, Keanu Sonn M. Fortaleza,  
Benny Gil A. Lactaotao, Marven Joffre G. Luis, Georcelle Leisley U. Nuarin,  
Sean Drei A. Octavo, and Gebreyl Isaac P. Rabang

Department of Computing and Information Studies, Saint Louis University

CS421, CS THESIS 2

Josephine S. Dela Cruz, DIT

May 15, 2026

## Abstract

Pulmonary Tuberculosis (PTB) remains a critical public health challenge in the Cordillera Administrative Region (CAR), Philippines, where complex patient profiles and programmatic gaps contribute to persistent treatment failure. This study develops machine learning models to predict treatment success or failure among PTB patients in DOH-CAR TB-DOTS facilities covering the period 2015–2025. A standardized preprocessing pipeline performs schema harmonization, patient pseudonymization, and a diagnostic-first missing data handling protocol that classifies each variable’s missingness mechanism—Missing Completely at Random (MCAR), Missing at Random (MAR), or Missing Not at Random (MNAR)—before routing it to one of four targeted strategies: column exclusion, listwise deletion, missing indicator imputation, or stochastic Multivariate Imputation by Chained Equations (MICE). Supervised ensemble tree-based classifiers—XGBoost, LightGBM, and Random Forest—are trained alongside Recurrent Neural Network (RNN) and Long Short-Term Memory (LSTM) temporal architectures to capture longitudinal treatment dynamics across the 6- to 12-month treatment course. Shapley Additive Explanations (SHAP) ensure model interpretability by quantifying the contribution of individual patient features to each prediction. The validated models are embedded in a Clinical Decision Support System (CDSS) that provides dynamic risk estimates and factor-level explanations to support early clinical intervention for patients at risk of treatment failure.

*Keywords:* tuberculosis, machine learning, treatment outcome prediction, clinical decision support system, explainable AI

**Treatment Success and Failure Prediction Using Machine Learning Models  
for Pulmonary Tuberculosis Patients Under the Department of Health-Cordillera  
Administrative Region (DOH-CAR) Tuberculosis-Directly Observed Treatment,  
Short-course (TB-DOTS) Program**

Pulmonary Tuberculosis (PTB) is a chronic infectious disease caused by *Mycobacterium tuberculosis* (MTB), a slow-growing and acid-fast bacillus that continues to challenge public health due to its gradual progression and risk of treatment failure (Tobin & Tristram, 2024). Despite effective therapies, PTB remains widespread due to its strong association with poverty, undernutrition, and comorbidities such as diabetes and HIV which weaken immune defenses and increase susceptibility (Tobin & Tristram, 2024). These factors delay diagnosis, complicate treatment, and contribute to disease progression. Globally, TB is still one of the leading causes of death, with 1.5 million deaths and 10 million new cases recorded in 2018, including half a million with rifampicin resistance.

In the Philippines, TB continues to be a major public health concern. The World Health Organization (WHO) identifies the country as one of the top contributors to the global TB burden as it accounts for an estimated 6.8 percent of cases. Before the COVID-19 pandemic, the Philippines had one of the highest TB rates in Asia with over 554 cases per 100,000 people in 2019 (Florentino et al., 2022). In 2023, around 37,000 Filipinos died from TB out of the 739,000 diagnosed cases (World Health Organization, 2023). The WHO classifies the country as a high burden for MDR-TB and TB/HIV co-infection (The Borgen Project, 2025). That same year, the DOH recorded 575,770 TB cases through its Integrated TB Information System which indicates that many cases remain unregistered or undiagnosed (Baron, 2025). The DOH implements CDC-recommended regimens for active and latent TB through the TB-DOTS strategy which was expanded in 2003 to include private clinics, pharmacies, and hospitals (Bendre et al., 2021; Department of Health [DOH], 2014). Despite these initiatives, gaps in diagnosis, treatment monitoring, and patient support persist especially in the management of MDR-TB (Stop TB Partnership, 2021).

In Baguio City, the DOH-CAR reported 437 confirmed TB cases in 2022 along with

over 4,405 missing or unreported cases making the area a priority for surveillance and treatment improvement (Department of Health - Republic of the Philippines, 2022), a gap driven by passive case-finding limitations, stigma, and financial barriers that allow undetected patients to continue transmitting the disease within their communities (Kaur & Sharma, 2021).

Contributing factors include high population density, internal migration, the city's cool and humid highland climate conducive to *Mycobacterium tuberculosis* transmission, and its role as a tertiary referral center that draws complex cases from neighboring provinces and municipalities, further amplifying the local disease burden (Khalid et al., 2022). Patient comorbidities such as diabetes mellitus, HIV infection, and malnutrition impair immune response and reduce treatment completion rates, while poverty-related barriers including food insecurity, substandard housing, and transportation costs disrupt adherence to directly observed treatment short-course (DOTS), compounded by health system limitations such as workforce shortages, inconsistent drug supply chains, and fragmented referral networks (Austin et al., 2021; Dong & Peng, 2013; Khalid et al., 2022; Sisk et al., 2023).

These challenges collectively highlight the need for locally adaptive, multi-pronged strategies, as treatment failure not only worsens individual patient outcomes but also prolongs community transmission and accelerates the emergence of multidrug-resistant tuberculosis (MDR-TB)—defined as resistance to at least rifampicin and isoniazid—which demands longer, more toxic, and costlier regimens with significantly lower success rates, posing a broader threat to regional disease control efforts in the Philippines (Aracri et al., 2025; Groenwold et al., 2012; World Health Organization, 2022).

### **Risk Factors for Treatment Failure**

To ensure uniformity in the evaluation of TB control efforts, the DOH has established standardized criteria for classifying treatment outcomes through the National Tuberculosis Program (NTP). Establishing these outcomes is crucial for monitoring program performance, ensuring

comparability across studies, and informing risk factor analysis. Ultimately, these outcomes of success or failure are not random but are linked to a distinct set of patient-specific, clinical, and programmatic risk factors. While mortality remains the primary outcome for most studies on treatment failure, some also consider the development of disease, resistance to multiple anti-TB drugs, and failure to follow up (Department of Health [DOH], 2014).

One of the risk factors is that TB patients have close contact with other patients with active TB infection. A study noted that TB patients who acquired the infection through close contact were shown to be at higher risk of delayed diagnosis and more likely to be non-adherent with the treatment either by refusing or discontinuing the regimen (Lamsal et al., 2009). Another risk factor is recurrence of the infection or having a history of receiving TB treatment. TB patients who missed follow-ups had a higher risk of reactivation of the infection and mortality as they are more difficult to treat compared to patients with no prior history of TB (Tok et al., 2020). Furthermore, missed or interrupted treatment causes the development of resistance to the drug if done frequently (Ndambuki et al., 2020).

Comorbidities are also discussed in studies that have consistently shown a significant influence on tuberculosis treatment outcomes across different settings. In Eastern Ethiopia, HIV co-infection was determined to be a substantial risk factor for treatment failure due to the weakened immune system which leads to a higher risk for ineffective treatment (Amante & Abdosh, 2015). In addition, TB and HIV drug interactions accelerate each other's progression and increase the occurrence of side effects or adverse effects which may compromise treatment success (Azeez et al., 2018). Likewise, a study in Davao City, Philippines, revealed that there was a higher risk of treatment failure in comorbid patients with 58.82% of the 51 patients who experienced treatment failure having at least one comorbidity (Hinay et al., 2025). Another major determinant is Diabetes Mellitus (DM). TB patients with diagnosed DM have a higher risk of treatment failure, relapse, and death compared to non-diabetics (Noubiap et al., 2019). In a systematic review and meta-analysis, patients with DM have a higher risk of experiencing treatment failure for Drug-resistant (DR) and mortality for MDR-TB with an odds ratio of 1.56 and 1.57 respectively (Xu et al., 2023). In line with

this, another study states that patients suffering from both PTB and chronic conditions, particularly DM, had reduced rates of sputum conversion with a higher probability of poor treatment outcomes including death and progression to MDR-TB (Siddiqui et al., 2016).

Taken together, these findings reveal that TB treatment failure is driven not by any single factor but by the complex interplay of comorbidities, programmatic gaps, and socioeconomic vulnerabilities that conventional single-factor analyses fail to capture, underscoring the need for an integrative machine learning framework capable of modeling these multidimensional interactions to enable earlier and more targeted clinical intervention (Amante & Abdosh, 2015; Department of Health - Republic of the Philippines, 2022; Hinay et al., 2025; Khalid et al., 2022; Noubiap et al., 2019; Peng et al., 2024; Rajpurkar et al., 2022; World Health Organization, 2022; Xu et al., 2023).

### **Machine Learning for TB Outcome Prediction**

Previous studies have examined tuberculosis treatment outcomes by focusing on sociodemographic characteristics, comorbidities, medical history, and other patient-level risk factors. However, the findings remain inconsistent and fragmented because methodological constraints limit how these factors are analyzed together. Most studies rely on traditional statistical analyses, such as statistical tests or regression models, that examine one factor at a time or a limited set of factors while assuming linear relationships among variables (Amante & Abdosh, 2015). These approaches provide important evidence on individual risk factors, but they may overlook the combined and non-linear effects that shape treatment success or failure.

Machine learning (ML) offers a distinct advantage in modeling complex and non-linear interactions among demographic, clinical, diagnostic, and treatment-related variables that conventional statistical analyses may overlook. By leveraging algorithms capable of learning from high-dimensional data, such as logistic regression, Support Vector Machines (SVM), and ensemble-based methods including Extreme Gradient Boosting (XGBoost), ML can uncover subtle dependencies that contribute to treatment success or failure (Rajpurkar et al., 2022). Previous research has demon-

strated the utility of ML in TB research, including predicting drug resistance, identifying high-risk patients, and optimizing treatment strategies (Kaur & Sharma, 2021; Kossakov et al., 2024). These applications show that predictive models can extend risk factor analysis by integrating multiple dimensions of patient data.

Beyond prediction, interpretability is necessary for applying ML in clinical contexts. Shapley Additive Explanations (SHAP), a post-hoc interpretability technique, quantifies the contribution of each variable to the model's predictions and enables a transparent evaluation of which patient features most strongly drive treatment success or failure. This supports explainable and trustworthy Artificial Intelligence (AI) in clinical settings (Peng et al., 2024). Using supervised learning with interpretable AI methods, a predictive framework can not only forecast treatment outcomes but also reveal the relative importance of risk factors such as drug resistance, comorbidities, or regimen adherence.

Taken together, the limitations of fragmented risk factor studies and the capabilities of interpretable ML highlight the need for an integrative approach that combines multiple dimensions of data in analyzing the interrelated variables that influence TB treatment outcomes. To address this gap, the present study adopts a holistic multi-layer analytical framework that integrates descriptive analysis and supervised predictive modeling using demographic, clinical, diagnostic, and treatment-related variables. Descriptive analyses are employed to characterize patient profiles and treatment patterns, while supervised learning models are trained directly on integrated patient data to predict treatment outcomes. This approach enables the models to capture individual-level patterns that improve predictive performance while maintaining interpretability (Shickel et al., 2018; Smith et al., 2021). The resulting models are designed to help healthcare providers identify patients at a higher risk of treatment failure, thereby enabling early intervention, tailored regimens, and improved patient management.

## Contributions

The DOH-CAR TB-DOTS program follows the standardized National Tuberculosis Program (NTP) protocol, wherein patient profiling begins at the facility level through passive case finding, where presumptive TB patients are identified based on clinical presentation such as chronic cough, weight loss, and fever, and subsequently subjected to diagnostic workup including sputum smear microscopy, GeneXpert MTB/RIF assay, and chest radiography (Bendre et al., 2021; Department of Health [DOH], 2014). Once diagnosed, patients are registered in the Treatment Card and Master List, assigned a registration group, and enrolled under a healthcare worker or trained treatment partner who administers and directly observes each dose throughout the six- to twelve-month regimen (Department of Health [DOH], 2014). Monthly follow-up visits are conducted to monitor sputum conversion, document vital signs and weight, and assess adherence, with outcomes eventually classified according to NTP criteria as Cured, Treatment Completed, Failed, Died, or Lost to Follow-up (Ting, 2011). However, this process is largely paper-based and retrospective in its risk assessment, meaning that clinical deterioration, non-adherence, and comorbidity-related complications are typically identified only after they have already influenced treatment trajectory rather than being anticipated and preemptively addressed, creating a critical gap wherein high-risk patients are not systematically flagged for early intervention, resulting in avoidable treatment failures, increased mortality, and continued MDR-TB emergence within the region (Khalid et al., 2022; World Health Organization, 2022).

To address these gaps, the main contributions of this paper are as follows. First, the study develops and validates an Integrated Predictive Framework comprising machine learning models trained on demographic, clinical, diagnostic, and treatment-related variables to accurately predict therapy success or failure among PTB patients in the CAR, thereby supporting early identification of at-risk individuals and providing evidence-based tools for improving TB program outcomes. Second, the study contributes a reproducible Data Preprocessing and Harmonization pipeline that standardizes and documents patient records from TB-DOTS facilities across the CAR through schema harmonization, patient pseudonymization, error correction, and a diagnostic-first missing



data handling protocol, ensuring dataset readiness for machine learning applications. Third, a Descriptive and Exploratory Analysis is conducted to characterize the demographic, clinical, diagnostic, and treatment profiles of the study population, quantify variable associations, and identify key predictors of treatment outcomes over the course of therapy. Fourth, the study incorporates Temporal Modeling through RNN and LSTM architectures that analyze time-dependent patterns in monthly patient monitoring data, enabling the detection of dynamic changes throughout the treatment course and improving the accuracy of outcome predictions to support timely clinical interventions. Fifth, the validated models are embedded within a Clinical Decision Support System that visualizes patient risk levels, highlights contributing factors via SHAP-based explanations, and dynamically updates risk estimates as new follow-up data become available, enabling healthcare providers to intervene earlier and more precisely for patients at risk of poor outcomes.

## **Research Objectives**

This study aims to develop and validate machine learning models that accurately predict treatment success or failure among TB patients in the CAR, Philippines. The integration of demographic, clinical, and laboratory variables focuses on supporting early identification of patients at risk of unsuccessful treatment and providing evidence-based tools for improving TB program outcomes, with the overarching goal of enabling personalized clinical interventions and efficient allocation of healthcare resources. Specifically, the research aims to achieve the following objectives:

1. To document and pre-process patient records of PTB cases from TB-DOTS facilities in the CAR through data cleaning, feature engineering, and reporting for readiness to be used in ML applications.
2. To describe the demographic, clinical, diagnostic, and treatment characteristics of PTB patients in the dataset, providing an overview of the study population.
3. To identify contributing variables that influence treatment outcomes over the course of

therapy, enabling recognition of the most critical factors that affect patient recovery and guiding more precise treatment adjustments that can be intervened with.

4. To incorporate temporal modeling that analyzes time-dependent patterns in patient data, allowing the detection of dynamic changes and trends throughout treatment, thereby improving the accuracy of outcome predictions and supporting timely clinical interventions.
5. To train and validate supervised learning models for predicting treatment success or failure among PTB patients, enabling early identification of individuals at risk of poor outcomes and ensuring that the model maintains high predictive accuracy and reliability when applied to evaluate incoming TB cases.
6. To develop a decision support tool that visualizes patient risk levels and predictive outputs.

## **Scope and Limitations**

The dataset will include PTB patients for the period 2015–2025 in the CAR specifically Baguio City excluding those who are still undergoing treatment at the time of data extraction. Furthermore, only patients with a well-defined treatment outcome of either success or failure will be included. The data gathered will encompass demographic information, laboratory tests, diagnostic data, clinical data, treatment regimen, follow-up examinations, drug administration data, and treatment outcomes. The study faces limitations related to the sufficiency and quality of the dataset. The data sourced from existing records may present challenges related to completeness, consistency, and representativeness due to limited access to complete clinical parameters. This could affect the inclusion of important risk factors and weaken the robustness of the analysis. Additionally, the study may not fully represent the current TB patient population as patients currently undergoing treatment are excluded. Difficulties are also expected during data extraction from DOH-CAR as variations in record-keeping over the 10-year period from 2015 to 2025 may be inconsistent due to changes in the healthcare system and shifts in documentation standards. Reliance on self-reported indicators may also introduce bias and inconsistencies that could affect

the predictive capabilities of the model.

### **Significance of the Study**

The resulting models are designed to help healthcare providers identify patients at a higher risk of treatment failure thus enabling early intervention, tailored regimens, and improved patient management. Ultimately, this machine learning framework supports the broader objective of strengthening TB control efforts in the Cordillera Administrative Region (CAR) by transforming existing patient data into predictive, explainable, and actionable tools for evidence-based public health practice. The results will help improve patient management and treatment planning, strengthen program implementation, and guide public health decision-making to improve the TB treatment success rate in CAR. Furthermore, the study aims to maximize benefits for the public by informing health administrators about the model which may be further developed into a solution or application to aid the local TB-DOTS program. The processed and pseudonymized derivative of the DOH-CAR PTB Treatment Outcome Dataset produced by the study will also be published to support the advancement of clinical, epidemiological, and machine learning research.


### **Method**

This study uses a dataset of records of PTB patients from the Cordillera Administrative Region (CAR), Philippines, covering 2015–2025 and consolidated by the Department of Health–CAR (DOH-CAR). After ethical approval, records are pseudonymized and processed using a standardized preprocessing pipeline, including schema harmonization, error correction, and feature scaling, followed by a diagnostic-first missing data handling protocol described in the Missing Data Handling subsection below. Exploratory and cluster analyses are conducted to characterize the cohort and derive subgroup-level features. Supervised ensemble tree-based classifiers, specifically XGBoost, LightGBM, and Random Forest, are trained to predict treatment outcomes defined by

National Tuberculosis Program criteria. To capture longitudinal treatment dynamics, temporal models using Recurrent Neural Networks (RNN) and Long Short-Term Memory (LSTM) are evaluated alongside static models, with ensemble strategies applied for improved robustness. Model performance is assessed using standard classification metrics, and Shapley Additive Explanations (SHAP) are employed to ensure interpretability. The final models are integrated into a clinical decision support framework to provide dynamic risk estimates throughout the treatment course. An overview can be found in Figure 1.

**Figure 1**

*Data Management and Analysis Process*



## Data Acquisition and Preprocessing

The study utilizes a retrospective dataset of PTB patients from the Cordillera Administrative Region (CAR), Philippines, covering the period from 2015 to 2025. Data is centrally consolidated by the Department of Health-CAR (DOH-CAR), which manages reports from government TB-DOTS facilities and private practitioners under the Public-Private Mix DOTS (PPMD) scheme (Department of Health [DOH], 2014).

The raw dataset comprises demographic variables (e.g., age, sex, civil status), clinical indicators (e.g., BMI, blood pressure), and diagnostic results (e.g., GeneXpert, Smear Microscopy). To ensure data quality for machine learning applications, a reproducible preprocessing pipeline is implemented. This includes schema harmonization to standardize variable formats and units, and the removal of direct identifiers to strictly enforce patient privacy. Missing data, a common challenge in clinical registries, is addressed through a diagnostic-first protocol described in the Missing Data Handling subsection below.

In addition to schema harmonization and missing data handling, other critical pre-processing activities include data annotation and outlier management. Data annotation is systematically performed to explicitly define the target variable for the predictive models. Based on the standardized National Tuberculosis Program (NTP) criteria, treatment outcomes are consolidated and binarized: patients classified as ‘Cured’ or ‘Treatment Completed’ are annotated as ‘Success’, whereas those marked as ‘Failed’, ‘Died’, or ‘Lost to Follow-up’ are annotated as ‘Failure’. Furthermore, extreme outliers in continuous clinical variables are evaluated and capped to prevent model distortion during training.

Crucially, to ensure the integrity of the predictive models and rigorously prevent data leakage, the dataset is partitioned into training and testing sets prior to the application of any data transformation steps. A stratified train-test split is performed to preserve the baseline distribution of treatment success and failure outcomes across both partitions. All subsequent preprocessing mechanisms—most notably feature scaling, normalization, and the multivariate imputation of missing values—are fitted exclusively on the training set. The derived parameters (such as the

mean and standard deviation) and imputation rules are then transformed onto the held-out test set. This strict isolation guarantees that the models do not inadvertently learn from the statistical distribution of the testing data, thereby providing a robust, realistic evaluation of their predictive performance on unseen patient records.

## Exploratory Data Analysis

Prior to predictive modeling, an exploratory analysis was conducted to characterize the study population, assess data completeness, and quantify the associations among variables to inform downstream preprocessing and modeling decisions.

**Table 1**

*Missing Data Summary by Column*

| Column                      | Missing_Count | Missing_Percentage |
|-----------------------------|---------------|--------------------|
| Days_To_Microscopy_Result   | 7297          | 98.75              |
| Days_To_RDT_Result          | 6638          | 89.84              |
| Days_Screening_To_Diagnosis | 4444          | 60.14              |
| Days_To_Treatment           | 4378          | 59.25              |
| Diagnosis_Season            | 3953          | 53.5               |
| Diagnosis_DayOfYear         | 3953          | 53.5               |
| Diagnosis_Quarter           | 3953          | 53.5               |
| Diagnosis_Month             | 3953          | 53.5               |
| Microscopy Result           | 2792          | 37.79              |
| Brgy                        | 1026          | 13.89              |
| City/Municipality           | 105           | 1.42               |
| Province                    | 71            | 0.96               |
| Treatment Health Facility   | 29            | 0.39               |

| Column                               | Missing_Count | Missing_Percentage |
|--------------------------------------|---------------|--------------------|
| Region                               | 15            | 0.2                |
| Date of Outcome/Status               | 2             | 0.03               |
| Bacteriologic Status                 | 1             | 0.01               |
| Type                                 | 0             | 0                  |
| Validation Status                    | 0             | 0                  |
| Age                                  | 0             | 0                  |
| Sex                                  | 0             | 0                  |
| Anatomical Site                      | 0             | 0                  |
| Source of Patient                    | 0             | 0                  |
| Screening/Diagnosing Health Facility | 0             | 0                  |
| RDT Result                           | 0             | 0                  |
| Year                                 | 0             | 0                  |
| Outcome/Status                       | 0             | 0                  |
| Registration Group                   | 0             | 0                  |
| Computed_Age                         | 0             | 0                  |

A global missingness audit was performed across all 28 columns to identify the sparsity structure of the dataset (Table 1). Per-column missing rates were ranked, revealing that variables such as microscopy result (37.8% missing), diagnosis-derived temporal features (53.6%), and clinical delay intervals (59.3–98.8%) exhibited substantial missingness. This audit provided the empirical basis for the missingness thresholds applied in the Missing Data Handling subsection.

Pairwise correlations among continuous variables were computed using Pearson’s  $r$  on available-case observations. The numerical set included age, four derived clinical delay intervals (days from screening to diagnosis, days to treatment, days to microscopy result, and days to RDT result), and diagnosis day-of-year. The purpose was to determine whether multivariate imputation via MICE was warranted: if continuous predictors carry meaningful correlation structure, MICE can borrow signal from observed features to estimate missing values; if uncorrelated, univariate mean-fill would be equally effective.

For categorical variables, pairwise association strength was quantified using Cramér’s V, a chi-square-based measure normalized to the  $[0, 1]$  interval. Columns with fewer than two unique non-null values or fewer than ten non-null observations were excluded from the matrix. The categorical set examined included patient demographics (sex, type of TB, anatomical site), and clinical classifications (registration group, bacteriologic status, and microscopy results).

Furthermore, to identify latent subgroups within the study population, the exploratory phase employs cluster analysis. Specifically, algorithms such as K-prototypes are utilized to effectively handle the mixture of categorical and continuous variables across demographic, clinical, and diagnostic dimensions. The optimal number of clusters is determined iteratively using evaluation metrics such as the Silhouette score and the elbow method to detect robust, natural groupings of patients. These clusters reveal distinguishing features such as specific comorbidity profiles, adherence behaviors, and regimen responses.

Importantly, the resulting cluster assignments are not merely descriptive; they are integrated as engineered features into the subsequent supervised learning models. To operationalize this, the assigned cluster labels are one-hot encoded to serve as supplementary categorical input variables for the tree-based classifiers (XGBoost, LightGBM, and Random Forest) and incorporated into the feature space for the temporal neural networks (RNN and LSTM). This approach enables the algorithms to learn both individual-level attributes and subgroup-level patterns that significantly influence treatment success or failure.

Additionally, the study employs cluster analysis on demographic, clinical, and diagnostic dimensions to detect natural groupings of patients. These clusters reveal distinguishing features such as specific comorbidity profiles, adherence behaviors, and regimen responses. Importantly, the resulting cluster assignments are not merely descriptive; they are integrated as engineered features into the subsequent supervised learning models. This approach enables the algorithms to learn both individual-level attributes and subgroup-level patterns that significantly influence treatment success or failure.

To assess data quality of the temporal monitoring dataset prior to preprocessing, a global missingness audit was conducted across all features. Missing rates were computed per column and ranked in descending order. Features exceeding 80% missingness were flagged for the column exclusion pathway, while features between 50% and 80% missingness were flagged for the missing indicator with fill pathway, contingent on predictive importance. This threshold-based triage follows established missing data handling frameworks and ensures that handling strategies are matched to the severity of missingness rather than applied uniformly



(Emmanuel et al., 2021).

Patient retention across the temporal dataset’s treatment months was analyzed by computing the proportion of patients with non-missing adherence data at each monthly interval from M0 through M12. Because retention declines naturally over a six- to twelve-month treatment course, a simple per-month count is insufficient; instead, the decay in available observations was characterized as a percentage of the initial cohort at each successive month. This analysis directly informs the temporal modeling scope: months beyond which fewer than 25% of patients retained data are excluded from feature-level correlation analyses to avoid estimates driven by a non-representative surviving subset.

Pairwise correlations among vital signs in the temporal monitoring dataset were computed using Pearson’s  $r$  on available-case observations. The presence and strength of inter-vital correlations directly supports the choice of multivariate imputation over univariate strategies: if vital signs are meaningfully correlated, a multivariate method such as Multiple Imputation by Chained Equations (MICE) can borrow signal from observed features to estimate missing values; if uncorrelated, MICE offers no advantage over mean- or median-fill (Aracri et al., 2025). The vital set examined included age, heart rate, oxygen saturation, systolic blood pressure, diastolic blood pressure, weight, height, respiratory rate, and temperature, where present in the temporal records.

To examine how temporal treatment features relate to each other, Pearson correlations were computed on per-patient observed means across Months 0 through 7. For each temporal feature — including dose adherence metrics (monthly doses taken, cumulative doses taken, percentage adherence), anthropometric measurements (weight, height), and diagnostic results (smear microscopy, Xpert MTB/RIF) — the observed values across a patient’s available monthly visits were averaged. Retention characteristics (last month attended, total observed months) were derived from the same window. This approach preserves the dropout signal: a patient who stops attending after M4 contributes means based only on M0–M4, and their limited observation window itself becomes informative rather than requiring imputation.

For categorical variables within the temporal dataset, pairwise association strength was quantified using Cramér’s  $V$ , a chi-square-based measure normalized to the  $[0, 1]$  interval. A mixed Pearson-Cramér’s matrix was not used, as Pearson’s  $r$  assumes interval-level measurement and produces misleading values when applied to nominal or ordinal categories with unequal category spacing (Kossakov et al., 2024). Computing Cramér’s  $V$  separately on categorical features avoids this assumption violation and produces


interpretable association estimates. Columns with fewer than two unique non-null values or fewer than ten non-null observations were excluded from the matrix, as association cannot be meaningfully estimated under those conditions.

To identify the extent of inconsistent coding that necessitated categorical harmonization, a categorical cardinality audit was conducted across all categorical features in the raw dataset. For each feature, the number of unique raw values was counted to quantify the degree of entry inconsistency. Features with cardinality far exceeding their expected number of valid categories — such as clinical observation fields that should contain a small set of standardized codes but instead contained hundreds of free-text variants — were flagged for harmonization in Stage 4 of the preprocessing pipeline. This audit directly motivates the harmonization targets by providing an empirical baseline of the data entry inconsistency problem before any cleaning was applied.

Finally, the temporal stability of adherence behavior across the treatment course was examined through lag-1 Pearson correlations of percentage adherence across consecutive month pairs (M0–M1 through M7–M8). Month-to-month correlation strength determines whether a recurrent architecture, which assumes sequential dependence between time steps, is warranted. If consecutive months are weakly correlated, a static model operating on aggregated patient means would be equally informative; if strongly correlated, the temporal structure justifies the additional complexity of an RNN or LSTM (Laine & Aila, 2016; Nayebi et al., 2022).

## **Missing Data Handling**

Clinical TB registries exhibit heterogeneous missingness arising from inconsistent documentation, varying facility capacities, and the longitudinal structure of treatment monitoring. Applying a single imputation strategy uniformly across all variables is statistically inappropriate when the mechanism generating missingness differs across features. Accordingly, this study implements a diagnostic-first missing data pipeline that first classifies each variable’s missingness mechanism and then routes the variable to one of four handling strategies: column exclusion, listwise deletion, missing indicator with fill, or stochastic Multiple Imputation by Chained Equations (MICE) imputation (Figure 2) (Dong & Peng, 2013; Emmanuel et al., 2021).

**Figure 2***Missing Data Handling Decision Logic*

The first diagnostic step applies Little's MCAR test at the dataset level to evaluate whether data are Missing Completely at Random (MCAR). Under the null hypothesis of MCAR, missingness is independent of all observed and unobserved values, meaning that complete cases constitute an unbiased subsample of the population and listwise deletion would be permissible. The test returned a statistically significant result ( $p < 0.05$ ), rejecting MCAR and confirming that missingness patterns in the dataset are systematically

related to the observed data, thus ruling out blanket listwise deletion (Dong & Peng, 2013).

Because Little’s test yields a single omnibus verdict at the dataset level, per-column mechanism classification requires a separate procedure. Following the dataset-level assessment, each variable’s missingness indicator is examined against the observed data to assign a column-level mechanism. For temporal (monthly) variables, Spearman rank correlations are computed between the column’s monthly missingness rate and treatment month; a monotonically increasing correlation indicates a dropout pattern classified as Missing Not at Random (MNAR). For static baseline variables, point-biserial correlations between the binary is-missing indicator and observed features are evaluated; significant associations suggest that missingness is explained by other observed variables (Missing at Random, or MAR), whereas non-significant associations indicate MNAR. This per-column classification, rather than a single dataset-level verdict, allows the pipeline to match each handling strategy to the specific statistical properties of each variable (Dong & Peng, 2013; Emmanuel et al., 2021).

To support the column exclusion and MICE eligibility decisions, a Random Forest classifier is trained on available cases, defined as rows where the treatment outcome is observed. Features are filled with column medians for ranking purposes only (available-case analysis), providing a coarse but sufficient signal for importance estimation. Variables in the bottom quartile of predictive importance become eligible for soft column exclusion, while variables above the median importance threshold are prioritized for stochastic imputation (Aracri et al., 2025).

Variables exceeding 80% missingness are excluded entirely under a hard threshold, regardless of their importance ranking, because reliable imputation is statistically untenable at such sparsity levels. A soft threshold additionally excludes variables with more than 50% missingness whose importance falls in the bottom quartile, on the grounds that the information gain from imputing a low-value sparse variable does not justify the synthetic variance introduced. Temporal variables are exempt from both thresholds, as their structural missingness reflects the longitudinal design of TB treatment monitoring rather than data quality problems (Junaid et al., 2025).

Listwise deletion is included as a pathway for variables confirmed MCAR at the column level with fewer than 15% missing observations, provided that the remaining sample size after deletion exceeds 500 patients. Under these conditions, the complete cases constitute an unbiased subsample and listwise deletion yields unbiased parameter estimates without the risk of variance distortion introduced by imputation. On the

present dataset, no variables satisfied all three conditions simultaneously, and this pathway was not activated (Dong & Peng, 2013).

Variables classified as MNAR and all temporal features, regardless of mechanism, are handled through a two-part strategy: first, a binary missingness indicator (`is_missing_[variable]`) is added as a new feature encoding the structural absence; second, the original variable is filled with the last observed value carried forward for temporal variables, or with the column median for continuous static variables and column mode for categorical static variables. This approach treats absence as informative rather than as noise to be imputed: in TB monitoring data, a missing Month 9 weight observation frequently signals treatment dropout, and the indicator captures this clinical pattern for the predictive model. Backward-fill is explicitly excluded from temporal imputation to prevent temporal data leakage, whereby a future measurement would contaminate a prior time step that the model cannot legitimately access at inference time (Austin et al., 2021; Groenwold et al., 2012; Sisk et al., 2023).

The four clinical vital signs confirmed MAR and ranking above the median importance threshold (oxygen saturation, heart rate, systolic blood pressure, diastolic blood pressure) are imputed using stochastic Multiple Imputation by Chained Equations with an ExtraTreesRegressor as the conditional model—a nonparametric, tree-based estimator that captures non-linear relationships in mixed-type clinical data without requiring variable standardization, analogous to the missForest framework (Aracri et al., 2025). This estimator has been shown to outperform linear and  $k$ -nearest neighbor alternatives in mixed-type clinical datasets. To address the fraction of missing information (FMI) for each variable, the number of MICE iterations is scaled as  $\max(20, \min(100, \lfloor \text{miss\_rate} \times 100 \rfloor))$ , ensuring sufficient convergence for high-sparsity features while bounding computational cost; this scaling follows an analogous principle to Von Hippel’s recommendation that the number of imputations should approximate the percentage of missing information (von Hippel, 2018), applied here to iteration counts rather than the number of imputed datasets. Stochastic rather than deterministic imputation is used to preserve the natural distributional variance of the imputed variables rather than collapsing them to conditional means (Alruhaymi & Kim, 2021).

## Identification of Variables Associated With Treatment Outcomes

To bridge the gap between predictive performance and clinical interpretability, the study identifies critical variables associated with treatment outcomes. While machine learning models are adept at handling high-dimensional data, understanding the “why” behind a prediction is crucial in medical contexts (Rajpurkar et al., 2022).

The framework utilizes Shapley Additive Explanations (SHAP), a post-hoc interpretability technique that quantifies the marginal contribution of each variable to the model’s output (Lundberg & Lee, 2017a). This allows for the precise identification of risk factors such as drug resistance, specific comorbidities, or delayed diagnosis that drive a patient toward “Success” (Cured, Completed) or “Failure” (Failed, Died, Lost to Follow-up; Peng et al., 2024). This method ensures the AI system remains transparent and supports evidence-based clinical decision-making (Kaur & Sharma, 2021; Kossakov et al., 2024).

## Temporal Modeling

Recognizing the dynamic nature of TB treatment, which spans 6 to 12 months, this study incorporates temporal modeling to analyze longitudinal patient monitoring data. Static models trained only on baseline features may fail to capture the evolution of adherence, follow-up attendance, and laboratory measurements throughout treatment. Accordingly, temporal models are evaluated to support month-by-month risk estimation from treatment Month 0 (M0) through Month 12 (M12).

Temporal modeling was implemented in two iterations, denoted as version 1 (v1) and version 2 (v2). Both iterations originate from the same DOH-CAR TB-DOTS temporal dataset lineage, but v2 represents a corrected and improved modeling-ready preparation motivated by methodological review of v1. In v1, an initial temporal pipeline was used as a baseline for modeling, including patient-level stratified splitting and progressive temporal expansion across M0–M12. However, subsequent review identified limitations that could undermine clinical defensibility, including increased risk of outcome leakage from late-treatment fields, inconsistent handling of imputed or unreliable outcome labels, and less rigorously audited cleaning and temporal feature construction.

Version 2 (v2) retrained the temporal models after correcting these issues. The v2 workflow uses

the audited cleaned temporal dataset output, filters supervised learning to patients with observed treatment outcomes, and explicitly removes leakage-prone and administrative identifier fields (e.g., outcome-date fields and late timeline variables not available at early inference). Temporal predictors are encoded in two forms: sequence inputs for deep temporal models and engineered temporal summaries for tree-based models. The deep temporal architecture uses a Hybrid Bidirectional LSTM (Bi-LSTM) to model long-term dependencies while mitigating vanishing gradients (Lu et al., 2022), and its recurrent structure is suited for month-to-month dependencies (Nayebi et al., 2022). For tree-based temporal baselines, monthly values are flattened and augmented with aggregates, trends, latest-observed values, and missingness indicators, and class imbalance is handled using training-only resampling (e.g., SMOTE-ENN). Finally, a hybrid framing is retained where temporal predictions can be compared alongside strong static learners (e.g., XGBoost) to improve robustness across heterogeneous patient subgroups (Laine & Aila, 2016).

## **Supervised Learning Model Development and Validation**

The core predictive component involves training and validating supervised learning models to classify treatment outcomes. The study leverages ensemble tree-based algorithms capable of modeling complex, non-linear interactions among mixed-type clinical variables while natively handling class imbalance: specifically XGBoost (Chen & Guestrin, 2016; Rajpurkar et al., 2022), LightGBM (Ke et al., 2017), and Random Forest (Breiman, 2001). These three classifiers were selected after a preliminary evaluation of six algorithms (which also included Logistic Regression, SVM, and Gradient Boosting); the excluded models were found to produce either poor discriminative performance ( $CV\text{-}AUC \leq 0.63$ ) or near-zero recall on the minority failure class, rendering them clinically unsafe for deployment in a CDSS. Two sampling strategies are compared for each classifier: no resampling (with native class-weighting) and SMOTE combined with Edited Nearest Neighbors (SMOTE-ENN), which increases the training-set failure class by 433% to a near-balanced 0.70:1 ratio.

### ***Feature Set Versions***

Three progressive feature set versions are evaluated to assess the incremental value of additional clinical and temporal features. The *baseline* set (11 features) includes core demographic and diagnostic variables: age, sex, anatomical site, bacteriologic status, microscopy result, registration group, source of patient, TB type, province, days to treatment, and year of registration. The *improved* set (14 features) augments the baseline with three facility-level geographic variables: city/municipality, treatment health facility, and screening/diagnosing health facility. The *extended* set (21 features) further incorporates seven temporal and diagnostic delay features derived from date engineering: diagnosis month, diagnosis quarter, diagnosis season (using the Philippine three-season meteorological calendar: *Tag-lamig* for cool dry December–February, *Tag-init* for hot dry March–May, and *Tag-ulan* for the rainy season June–November), days from screening to diagnosis, days from diagnosis to microscopy result, days from diagnosis to rapid diagnostic test (RDT) result, and RDT result (binned from 42 raw variants into three classes: Detected, Not Detected, and Not Done). The date of outcome was excluded from all feature sets to prevent causal leakage, as this field is generated downstream of the treatment result.

### ***Data Partitioning and Class Imbalance***

The dataset is partitioned using a stratified 70/20/10 patient-level split into training (70%), test (20%), and holdout validation (10%) sets, preserving class proportions across all subsets (Hastie et al., 2009). The severe class imbalance (approximately 86% success versus 14% failure) is addressed on the training set only using three oversampling strategies: SMOTE (Chawla et al., 2002), which generates synthetic minority-class observations via interpolation; SMOTE-ENN (Batista et al., 2004; Nishat et al., 2022), which applies SMOTE followed by Edited Nearest Neighbors undersampling to remove ambiguous borderline samples; and SMOTE-Tomek (Batista et al., 2004), which applies SMOTE followed by Tomek link removal to clean the class boundary. All resampling is applied exclusively to the training set, preserving unaltered class distributions in the test and validation sets for unbiased evaluation (Salmi et al., 2024).



### ***Cross-Validation and Model Selection***

Each model configuration is evaluated using five-fold stratified cross-validation on the training data to estimate generalization performance, yielding a mean and standard deviation of the area under the receiver operating characteristic curve (ROC-AUC) across folds (Hastie et al., 2009). Because the primary clinical objective is the early identification of patients at risk of treatment failure, model selection employs a composite scoring metric:  $\text{score} = 0.6 \times \text{AUC} + 0.4 \times \text{Recall}(\text{Failure})$ . This weighting prioritizes discrimination ability while explicitly rewarding sensitivity to the minority failure class. The best-scoring configuration across all feature versions and resampling strategies is then evaluated on the held-out 10% validation set to obtain an unbiased estimate of real-world performance.

To ensure reproducibility and eliminate transcription errors, all experiments are executed through a centralized pipeline that standardizes training, cross-validation, and evaluation, and automatically exports all performance metrics and visualizations directly to native  $\text{\LaTeX}$  format for inclusion in this manuscript.

### **Development of a Clinical Decision Support Tool**

The validated models are embedded into a Clinical Decision Support System (CDSS) designed for healthcare providers. This system visualizes patient risk levels and provides a summary of predictive outputs. Key features include a predictive dashboard that displays the probability of treatment failure and success, risk factor visualization that highlights the specific variables (e.g., adherence gaps, lab results) contributing to the risk score, and dynamic updates where the system accepts new follow-up data (e.g., monthly weight, smear results), triggering updated risk predictions to support real-time clinical management throughout the treatment duration.

### **SLM-Based Narrative Explanation Layer**

To improve the usability of model explanations for healthcare users, the CDSS includes a small language model (SLM) as a constrained narrative layer that verbalizes model outputs into readable clinical text. SHAP-based feature contribution outputs remain the primary explanation source, while the SLM

functions as a presentation layer rather than an independent predictor or diagnostic engine (Lundberg & Lee, 2017b; Zeng & Zhu, 2024). This component operates post-prediction and does not alter model training, inference probabilities, or risk classification thresholds. In the implemented pipeline, the web client constructs a structured interpretation request containing selected patient context fields, predicted treatment-failure probability, and ranked contribution signals, then sends the payload to a dedicated interpretation endpoint for streamed narrative rendering in the interface.

The SLM is operationally bounded to explanation tasks only. It is not used to infer new clinical causes, generate treatment prescriptions, or replace clinician judgment. This role separation aligns with current medical-domain SLM framing and health AI governance expectations for bounded decision support, while keeping model interpretability grounded in the original attribution signal (Collins et al., 2024; Sellergren et al., 2025; World Health Organization, 2021). Final clinical judgment remains with the treating healthcare professional.

## **Results and Discussion**

This section presents the results of machine learning analyses aimed at predicting treatment success and failure among PTB patients under the DOH-CAR TB-DOTS program. Descriptive statistics summarize patient demographics, clinical characteristics, and treatment outcomes, while predictive model performance is evaluated using test and validation datasets. Both static tree-based models (XGBoost, LightGBM, Random Forest) under two sampling conditions and temporal models are compared, and feature importance analyses provide insights into the key factors influencing treatment outcomes. The discussion interprets these findings in the context of clinical relevance and public health priorities, highlighting implications for early intervention, resource allocation, and strengthened TB management.

### **Missing Data Handling Outcomes**

Of the 44 features identified with missing data, 15 were excluded via column exclusion. All 15 exceeded the hard missingness threshold of 80%, including near-complete variables such as tuberculosis culture (99.8%), other laboratory tests (99.8%), and tuberculin skin test (93.5%). One additional variable,

the raw blood pressure string field, was soft-excluded at 67.3% missingness and below-median predictive importance, as it had already been parsed into the structured systolic and diastolic blood pressure variables. All 599 patient records were retained in the dataset, with zero row deletion. Full sample retention is particularly consequential given the small cohort size and the severe class imbalance of 86.2% treatment success versus 13.8% failure; listwise deletion under these conditions would have compounded the minority class representation problem (Dong & Peng, 2013).

The listwise deletion pathway was not activated, as the dataset-level MCAR test was rejected and no individual variable simultaneously satisfied all three conditions of confirmed MCAR mechanism, fewer than 15% missing observations, and a post-deletion sample size exceeding 500. Twenty-five variables were routed to the missing indicator with fill strategy, encompassing eight temporal monthly features (including weight, percentage adherence, cumulative doses taken, and monthly smear results) and 17 static baseline variables with confirmed MNAR patterns. Four variables — oxygen saturation, heart rate, systolic blood pressure, and diastolic blood pressure — were imputed via stochastic MICE following confirmation of a MAR mechanism; their missingness indicators were significantly associated with visit attendance under the point-biserial correlation procedure described above — a clinically plausible MAR mechanism, as vital sign recording depends on whether the patient attended the monthly monitoring visit rather than on the unobserved values of the measurements themselves. The concentration of stochastic imputation on this small subset of high-value features reflects the diagnostic-first design of the pipeline, which reserves computationally intensive probabilistic imputation for variables where the statistical assumptions required for its validity are met (Austin et al., 2021; Emmanuel et al., 2021).

**Table 2**

*Global Feature Missingness Rates in the Raw Temporal Dataset*

| Feature    | Missing Count | Missing % | Column Exclusion<br>( $\geq 80\%$ ) |
|------------|---------------|-----------|-------------------------------------|
| m8_weight  | 599           | 100.0     | Yes                                 |
| m12_weight | 599           | 100.0     | Yes                                 |

| Feature                                                   | Missing Count | Missing % | Column Exclusion<br>(≥80%) |
|-----------------------------------------------------------|---------------|-----------|----------------------------|
| m11_height                                                | 599           | 100.0     | Yes                        |
| m9_xpert_mtb_rif                                          | 599           | 100.0     | Yes                        |
| m9_smear_tb_lamp                                          | 599           | 100.0     | Yes                        |
| m7_xpert_mtb_rif                                          | 599           | 100.0     | Yes                        |
| m7_smear_tb_lamp                                          | 599           | 100.0     | Yes                        |
| m10_smear_tb_lamp                                         | 599           | 100.0     | Yes                        |
| m10_weight                                                | 599           | 100.0     | Yes                        |
| risk\_factor\_s\_for\_drug\_<br>_resistance\_tuberculosis | 599           | 100.0     | Yes                        |
| m12_height                                                | 599           | 100.0     | Yes                        |
| m12_smear_tb_lamp                                         | 599           | 100.0     | Yes                        |
| m12_xpert_mtb_rif                                         | 599           | 100.0     | Yes                        |
| m9_weight                                                 | 599           | 100.0     | Yes                        |
| m8_xpert_mtb_rif                                          | 599           | 100.0     | Yes                        |
| m8_smear_tb_lamp                                          | 599           | 100.0     | Yes                        |
| m1_xpert_mtb_rif                                          | 599           | 100.0     | Yes                        |
| m11_weight                                                | 599           | 100.0     | Yes                        |
| m10_xpert_mtb_rif                                         | 599           | 100.0     | Yes                        |
| m11_smear_tb_lamp                                         | 599           | 100.0     | Yes                        |
| m11_xpert_mtb_rif                                         | 599           | 100.0     | Yes                        |
| m3_xpert_mtb_rif                                          | 599           | 100.0     | Yes                        |
| m4_xpert_mtb_rif                                          | 598           | 99.83     | Yes                        |
| m10_height                                                | 598           | 99.83     | Yes                        |
| tuberculosis_culture                                      | 598           | 99.83     | Yes                        |
| m12_monthly_missed_doses                                  | 598           | 99.83     | Yes                        |
| other_lab_test                                            | 598           | 99.83     | Yes                        |

| Feature                    | Missing Count | Missing % | Column Exclusion<br>(>=80%) |
|----------------------------|---------------|-----------|-----------------------------|
| m12_pct_adherence          | 598           | 99.83     | Yes                         |
| m8_height                  | 598           | 99.83     | Yes                         |
| m9_height                  | 598           | 99.83     | Yes                         |
| m12_monthly_doses_taken    | 597           | 99.67     | Yes                         |
| m4_smear_tb_lamp           | 597           | 99.67     | Yes                         |
| m1_smear_tb_lamp           | 597           | 99.67     | Yes                         |
| others                     | 597           | 99.67     | Yes                         |
| m11_monthly_missed_doses   | 597           | 99.67     | Yes                         |
| m11_pct_adherence          | 597           | 99.67     | Yes                         |
| m12_cumulative_doses_taken | 597           | 99.67     | Yes                         |
| m10_monthly_missed_doses   | 596           | 99.5      | Yes                         |
| m9_pct_adherence           | 596           | 99.5      | Yes                         |
| m10_pct_adherence          | 596           | 99.5      | Yes                         |
| m11_monthly_doses_taken    | 595           | 99.33     | Yes                         |
| m9_monthly_missed_doses    | 595           | 99.33     | Yes                         |
| m11_cumulative_doses_taken | 595           | 99.33     | Yes                         |
| m8_pct_adherence           | 593           | 99.0      | Yes                         |
| m3_smear_tb_lamp           | 593           | 99.0      | Yes                         |
| m8_monthly_missed_doses    | 592           | 98.83     | Yes                         |
| m10_monthly_doses_taken    | 592           | 98.83     | Yes                         |
| m10_cumulative_doses_taken | 592           | 98.83     | Yes                         |
| m6_smear_tb_lamp           | 590           | 98.5      | Yes                         |
| m9_cumulative_doses_taken  | 588           | 98.16     | Yes                         |
| m7_height                  | 588           | 98.16     | Yes                         |
| m5_smear_tb_lamp           | 588           | 98.16     | Yes                         |
| m9_monthly_doses_taken     | 588           | 98.16     | Yes                         |

| Feature                   | Missing Count | Missing % | Column Exclusion<br>(≥80%) |
|---------------------------|---------------|-----------|----------------------------|
| m6_xpert_mtb_rif          | 587           | 98.0      | Yes                        |
| m2_smear_tb_lamp          | 587           | 98.0      | Yes                        |
| m0_smear_tb_lamp          | 586           | 97.83     | Yes                        |
| m7_weight                 | 586           | 97.83     | Yes                        |
| m5_xpert_mtb_rif          | 585           | 97.66     | Yes                        |
| m8_cumulative_doses_taken | 582           | 97.16     | Yes                        |
| m8_monthly_doses_taken    | 581           | 96.99     | Yes                        |
| dat_supported_dup         | 576           | 96.16     | Yes                        |
| m2_xpert_mtb_rif          | 572           | 95.49     | Yes                        |
| other                     | 570           | 95.16     | Yes                        |
| m5_height                 | 563           | 93.99     | Yes                        |
| m4_height                 | 562           | 93.82     | Yes                        |
| m6_height                 | 562           | 93.82     | Yes                        |
| m6_weight                 | 561           | 93.66     | Yes                        |
| tuberculin_skin_test      | 559           | 93.32     | Yes                        |
| m3_height                 | 558           | 93.16     | Yes                        |
| m5_weight                 | 557           | 92.99     | Yes                        |
| m2_height                 | 555           | 92.65     | Yes                        |
| m0_xpert_mtb_rif          | 554           | 92.49     | Yes                        |
| m4_weight                 | 554           | 92.49     | Yes                        |
| m1_height                 | 554           | 92.49     | Yes                        |
| nationality_raw           | 551           | 91.99     | Yes                        |
| m3_weight                 | 550           | 91.82     | Yes                        |
| dat_supported             | 549           | 91.65     | Yes                        |
| prior_history_of_tb       | 549           | 91.65     | Yes                        |
| m2_weight                 | 543           | 90.65     | Yes                        |

| Feature                                           | Missing Count | Missing % | Column Exclusion<br>( $\geq 80\%$ ) |
|---------------------------------------------------|---------------|-----------|-------------------------------------|
| m1_weight                                         | 540           | 90.15     | Yes                                 |
| m7_pct_adherence                                  | 519           | 86.64     | Yes                                 |
| regimen\_type\_at\_6th\_month\_of\_<br>_treatment | 518           | 86.48     | Yes                                 |
| smear_microscopy                                  | 502           | 83.81     | Yes                                 |
| m7_monthly_missed_doses                           | 502           | 83.81     | Yes                                 |
| respiratory_rate                                  | 501           | 83.64     | Yes                                 |
| height                                            | 492           | 82.14     | Yes                                 |
| temperature                                       | 481           | 80.3      | Yes                                 |
| weight                                            | 476           | 79.47     | No                                  |
| m7_cumulative_doses_taken                         | 465           | 77.63     | No                                  |
| o2_sat                                            | 462           | 77.13     | No                                  |
| m7_monthly_doses_taken                            | 453           | 75.63     | No                                  |
| heart_rate                                        | 441           | 73.62     | No                                  |
| blood_pressure                                    | 403           | 67.28     | No                                  |
| m6_pct_adherence                                  | 398           | 66.44     | No                                  |
| regimen\_type\_at\_end\_of\_<br>_treatment        | 355           | 59.27     | No                                  |
| m6_monthly_missed_doses                           | 347           | 57.93     | No                                  |
| m5_pct_adherence                                  | 336           | 56.09     | No                                  |
| m4_pct_adherence                                  | 319           | 53.26     | No                                  |
| m3_pct_adherence                                  | 312           | 52.09     | No                                  |
| m2_pct_adherence                                  | 302           | 50.42     | No                                  |
| m1_pct_adherence                                  | 293           | 48.91     | No                                  |
| name\_of\_treatment\_unit                         | 290           | 48.41     | No                                  |
| m0_pct_adherence                                  | 277           | 46.24     | No                                  |

| Feature                                   | Missing Count | Missing % | Column Exclusion<br>(≥80%) |
|-------------------------------------------|---------------|-----------|----------------------------|
| treatment\_regimen                        | 273           | 45.58     | No                         |
| m6_cumulative_doses_taken                 | 263           | 43.91     | No                         |
| co_morbidities                            | 253           | 42.24     | No                         |
| m5_monthly_missed_doses                   | 252           | 42.07     | No                         |
| m0_cumulative_doses_taken                 | 244           | 40.73     | No                         |
| continuation\_phase\_start\_date          | 239           | 39.9      | No                         |
| m3_monthly_missed_doses                   | 235           | 39.23     | No                         |
| m0_height                                 | 235           | 39.23     | No                         |
| date\_of\_outcome                         | 235           | 39.23     | No                         |
| m4_monthly_missed_doses                   | 233           | 38.9      | No                         |
| m6_monthly_doses_taken                    | 228           | 38.06     | No                         |
| continuation\_phase\_end\_date            | 228           | 38.06     | No                         |
| m0_monthly_missed_doses                   | 226           | 37.73     | No                         |
| m2_monthly_missed_doses                   | 222           | 37.06     | No                         |
| intensive\_phase\_end\_date               | 217           | 36.23     | No                         |
| m3_cumulative_doses_taken                 | 217           | 36.23     | No                         |
| drug\_resistance\_bacteriological\_status | 217           | 36.23     | No                         |
| chest\_x\_ray\_at\_case\_notification     | 213           | 35.56     | No                         |
| m1_monthly_missed_doses                   | 186           | 31.05     | No                         |
| m0_weight                                 | 185           | 30.88     | No                         |
| xpert_mtb_rif                             | 182           | 30.38     | No                         |
| civil_status                              | 167           | 27.88     | No                         |
| nationality                               | 166           | 27.71     | No                         |
| date\_of\_notification                    | 159           | 26.54     | No                         |



| Feature                                      | Missing Count | Missing % | Column Exclusion<br>(≥80%) |
|----------------------------------------------|---------------|-----------|----------------------------|
| m5_cumulative_doses_taken                    | 157           | 26.21     | No                         |
| m4_cumulative_doses_taken                    | 152           | 25.38     | No                         |
| name\_of\_diagnosing\_facility               | 150           | 25.04     | No                         |
| intensive\_phase\_start\_date                | 136           | 22.7      | No                         |
| outcome                                      | 127           | 21.2      | No                         |
| m2_cumulative_doses_taken                    | 126           | 21.04     | No                         |
| regimen\_type\_at\_start\_of\_<br>_treatment | 122           | 20.37     | No                         |
| case\_registration\_group                    | 117           | 19.53     | No                         |
| date\_of\_diagnosis                          | 110           | 18.36     | No                         |
| m5_monthly_doses_taken                       | 101           | 16.86     | No                         |
| m1_cumulative_doses_taken                    | 93            | 15.53     | No                         |
| m4_monthly_doses_taken                       | 90            | 15.03     | No                         |
| m3_monthly_doses_taken                       | 80            | 13.36     | No                         |
| bacteriologic\_status                        | 79            | 13.19     | No                         |
| diagnosis                                    | 62            | 10.35     | No                         |
| treatment\_start\_date                       | 57            | 9.52      | No                         |
| m2_monthly_doses_taken                       | 57            | 9.52      | No                         |
| m1_monthly_doses_taken                       | 39            | 6.51      | No                         |
| sex                                          | 33            | 5.51      | No                         |
| no                                           | 31            | 5.18      | No                         |
| m0_monthly_doses_taken                       | 30            | 5.01      | No                         |
| date\_of\_birth                              | 20            | 3.34      | No                         |
| age                                          | 18            | 3.01      | No                         |
| data\_year                                   | 0             | 0.0       | No                         |
| source\_file                                 | 0             | 0.0       | No                         |


| Feature  | Missing Count | Missing % | Column Exclusion<br>( $\geq 80\%$ ) |
|----------|---------------|-----------|-------------------------------------|
| facility | 0             | 0.0       | No                                  |

Table 2 presents the complete missingness profile of all features in the temporal dataset prior to preprocessing. Fifteen variables exceeded the 80% hard threshold and were routed to column exclusion, including tuberculosis culture (99.8%), other laboratory tests (99.8%), and tuberculin skin test (93.5%). An additional set of variables fell within the 50%–80% threshold band — including height (71.7%), weight (71.7%), smear microscopy (55.6%), and risk factor fields (91.7%) — where handling was contingent on predictive importance under the missing indicator with fill pathway.

**Table 3**

*Patient Data Retention by Month (M0–M12): Evidence for Missing Indicator with Fill*

| Month | Active Patients | % Retained | % Data Decay | Adherence Obs. |
|-------|-----------------|------------|--------------|----------------|
| M0    | 572             | 95.49      | 4.51         | 322            |
| M1    | 561             | 93.66      | 6.34         | 306            |
| M2    | 543             | 90.65      | 9.35         | 297            |
| M3    | 519             | 86.64      | 13.36        | 287            |
| M4    | 509             | 84.97      | 15.03        | 280            |
| M5    | 498             | 83.14      | 16.86        | 263            |
| M6    | 373             | 62.27      | 37.73        | 201            |
| M7    | 146             | 24.37      | 75.63        | 80             |
| M8    | 18              | 3.01       | 96.99        | 6              |
| M9    | 11              | 1.84       | 98.16        | 3              |
| M10   | 7               | 1.17       | 98.83        | 3              |
| M11   | 4               | 0.67       | 99.33        | 2              |
| M12   | 2               | 0.33       | 99.67        | 1              |

**Figure 3***Patient Retention Cascade Across Treatment Months*

Patient retention across the treatment course further clarifies the missingness mechanism. At treatment initiation (M0), 95.5% of patients had available data, declining to 83.1% at M5, then dropping sharply to 62.3% at M6 and 24.4% at M7. Beyond M7, fewer than 4% of patients remained observable. This monotonic decay confirms that missingness in the temporal features is driven by treatment dropout rather than random absence, directly supporting the missing indicator with fill pathway’s use of forward-fill with a missing indicator for temporal variables (Groenwold et al., 2012). The M7 cutoff also defines the temporal modeling scope: per-patient observed means are computed only across Months 0 through 7, as estimates beyond this window would rely on a non-representative surviving subset.

**Table 4**

*Descriptive Statistics and Skewness of Raw Vital Sign Features (Justifying Clinical Bound Clipping and MICE Imputation)*

| Feature | N   | Mean | Median | SD    | Min  | Max  | Skewness | Kurtosis | Missing % |
|---------|-----|------|--------|-------|------|------|----------|----------|-----------|
| age     | 581 | 46.2 | 48.0   | 21.24 | 0.42 | 93.0 | -0.06    | -0.824   | 3.01      |

*Table 4 continued*

| Feature          | N   | Mean  | Median | SD    | Min  | Max   | Skewness | Kurtosis | Missing % |
|------------------|-----|-------|--------|-------|------|-------|----------|----------|-----------|
| heart_rate       | 141 | 89.6  | 88.0   | 15.16 | 59.0 | 129.0 | 0.41     | -0.408   | 73.62     |
| respiratory_rate | 73  | 19.21 | 19.0   | 1.62  | 16.0 | 24.0  | 0.459    | 0.683    | 83.64     |
| temperature      | 91  | 36.39 | 36.4   | 0.26  | 35.6 | 37.0  | -0.501   | 0.555    | 80.3      |
| o2_sat           | 121 | 18.65 | 0.97   | 36.88 | 0.9  | 99.0  | 1.612    | 0.623    | 77.13     |


Table 4 reports the descriptive statistics and distributional properties of the vital sign features in the raw temporal dataset. The skewness values confirm several vitals exhibit moderate right-skew — particularly heart rate (skew = 1.52) and respiratory rate (skew = 1.83) — supporting the use of a tree-based MICE estimator (ExtraTreesRegressor) that does not assume normality. The four variables selected for stochastic MICE — oxygen saturation, heart rate, systolic blood pressure, and diastolic blood pressure — each show a clinically plausible MAR mechanism, with missingness linked to visit attendance rather than unobserved measurement values.


### **Dataset Characteristics and Class Distribution**

The final validated dataset comprised 7,389 complete patient records encompassing demographic, geographic, and clinical features. Overall, the dataset exhibited a significant class imbalance, with 86.2% of evaluated cases classified as treatment success and 13.8% as treatment failure. To ensure robust evaluation, the dataset was stratified into 70% training, 20% testing, and 10% validation subsets. Feature engineering expanded the dataset to 15 predictors, integrating spatial identifiers (Region, Province, City/Municipality) with temporal markers and derived clinical delays (e.g., Diagnosis-to-Treatment days).

**Figure 4**

*TB Case Notifications by Year (2015–2025)*




**Figure 5***Age-Sex Pyramid of the Study Cohort*

Resampling using SMOTE combined with Edited Nearest Neighbors (SMOTE+ENN) was applied exclusively to the training set. This technique increased the minority (failure) class by 433% to achieve a near-balanced 0.70:1 ratio, mitigating algorithm bias toward the majority class while preserving realistic, imbalanced evaluation conditions in the test and validation sets.

## **Outcome Distribution**

The initial distribution of treatment outcomes highlighted the programmatic successes as well as the persistent challenges within the DOH-CAR TB-DOTS program. The vast majority of the cohort achieved a positive outcome, with 65.1% of patients officially classified as “Treatment Completed” and 14.2% as “Cured.”

Conversely, adverse outcomes accounted for a critical minority of the dataset. Patients who were “Lost to Follow-up” made up 8.4% of the cohort, while 3.3% of the cases resulted in the patient having “Died.” A smaller fraction experienced outright treatment failure (0.9%). The remaining cases were either still “On Treatment” (6.9%) or “Not Evaluated” (0.9%) at the time of data consolidation.

**Figure 6***Distribution of Treatment Outcomes Among PTB Patients***Longitudinal Trends (2015–2025)**

An analysis of annual treatment outcomes from 2015 to 2025 revealed significant temporal shifts, reflecting both programmatic improvements and external systemic shocks:

1. **Treatment Success Rates:** The success rate remained relatively stable in the low-to-mid 80% range for most of the decade. Success rates dipped slightly to 81.03% in 2018 before steadily climbing to a peak of 88.92% in 2023.
2. **Mortality and Pandemic Impact:** The mortality rate among TB patients showed a distinct peak in 2020 (5.47%) and 2021 (6.63%), up from 3.65% in 2019. This sharp increase aligns with the severe disruptions caused by the COVID-19 pandemic, which delayed TB diagnoses, limited healthcare access, and exacerbated respiratory comorbidities.
3. **Loss to Follow-up:** The rate of patients lost to follow-up was highest between 2018 and 2020, peaking at 11.99%. However, this metric showed a steady and encouraging decline in subsequent years, dropping to 5.56% by 2024. This suggests improved patient tracking and adherence interventions by the local health program.
4. **2025 Cohort Anomaly:** The data for 2025 showed a drastic drop in recorded success (30.01%) and



a massive spike in unresolved cases (64.80%). This is indicative of an incomplete reporting year, where the majority of recent patients are still undergoing their 6- to 12-month treatment regimens and have not yet been assigned a final programmatic outcome.

**Table 5**

*Treatment Outcome Distribution by Year (%)*

| Year | Died | Lost to Follow-up | Other | Success |
|------|------|-------------------|-------|---------|
| 2015 | 2.65 | 9.56              | 3.82  | 83.97   |
| 2016 | 2.98 | 11.57             | 1.32  | 84.13   |
| 2017 | 3.84 | 8.8               | 1.28  | 86.08   |
| 2018 | 4.25 | 11.53             | 3.19  | 81.03   |
| 2019 | 3.65 | 11.8              | 3.51  | 81.04   |
| 2020 | 5.47 | 11.99             | 0.88  | 81.66   |
| 2021 | 6.63 | 6.82              | 0.97  | 85.58   |
| 2022 | 4.72 | 7.08              | 0.92  | 87.29   |
| 2023 | 3.92 | 6.35              | 0.81  | 88.92   |
| 2024 | 3.66 | 5.56              | 4.55  | 86.24   |
| 2025 | 2.05 | 3.14              | 64.8  | 30.01   |

## Diagnostic and Treatment Intervals

The temporal intervals between screening, diagnosis, and treatment initiation were extracted to identify operational delays, which are known risk factors for both disease transmission and treatment failure. Statistical summaries of these intervals revealed highly right-skewed distributions:

1. **Screening to Diagnosis:** The median time from initial screening to official diagnosis was 2.0 days. However, the upper quartile (Q3) extended to 21.09 days, indicating that while most patients are diagnosed rapidly, a subset experiences weeks of delay.
2. **Diagnosis to Treatment:** The median duration between diagnosis and treatment initiation was efficient at approximately 2.68 days. Similar to the diagnostic interval, however, the third quartile reached 13.12 days.
3. **Microscopy Results:** The median time to receive microscopy results was 6.60 days, with the third quartile extending to 14.95 days.


These skewed distributions demonstrate that while the median programmatic response is efficient, the long tails represent significant delays for high-risk patients. These interval metrics served as crucial continuous features in the machine learning models, allowing the algorithms to correlate logistical delays with the eventual probability of patient mortality or loss to follow-up.

**Table 6**


*Summary Statistics for Time Intervals (Days)*

| Interval                    | Count | Median | Q1 | Q3 |
|-----------------------------|-------|--------|----|----|
| Days_Screening_To_Diagnosis | 2945  | 0      | 0  | 7  |
| Days_To_Treatment           | 3011  | 1      | 0  | 6  |
| Days_To_Microscopy_Result   | 92    | 4      | 0  | 13 |
| Days_To_RDT_Result          | 751   | 0      | -1 | 3  |

## Correlation Analysis

**Figure 7***Pearson Correlation Matrix of Continuous Variables*

The Pearson correlation matrix (Figure 7) reveals the pairwise relationships among age, clinical delay intervals, and diagnosis timing. Age exhibited weak negative correlations with the time intervals ( $|r| < 0.10$ ), indicating that older patients experienced marginally shorter delays. The clinical delay variables showed moderate positive inter-correlations (screening-to-diagnosis with diagnosis-to-treatment:  $r = 0.32$ ), suggesting that delays at one stage of the care cascade are associated with delays at subsequent stages. Diagnosis day-of-year was essentially uncorrelated with all other variables ( $|r| < 0.05$ ), confirming that seasonal timing is independent of patient demographic and operational factors.

**Figure 8***Cramér's V Correlation Matrix of Categorical Variables*

The Cramér's V matrix (Figure 8) quantifies pairwise associations among the categorical feature set. Most pairwise associations fell in the moderate range ( $V \approx 0.3\text{--}0.5$ ), indicating meaningful but not redundant categorical structure. Geographic variables (province, region, and health facility) formed a cluster of elevated mutual associations ( $V > 0.6$ ), as expected given the hierarchical nesting of administrative units. The outcome variable showed modest associations with registration group and anatomical site ( $V \approx 0.25$ ), suggesting that categorical features carry a measurable but limited signal for treatment outcome prediction.

## Validation Set Performance and Generalization

**Table 7**

*Comprehensive Performance Metrics for All Static Model Configurations (Improved Feature Set)*


| Model          | Sampler          | AUC           | CV-AUC        | $\pm$ Std     | Acc.          | Prec. (F)     | Recall (F)    | F1 (F)        | Prec. (S)     | Recall (S)    | F1 (S)        | Score         |
|----------------|------------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|
| XGBoost        | None             | 0.7117        | 0.7024        | 0.0085        | 0.7065        | 0.2510        | 0.5656        | 0.3477        | 0.9127        | 0.7291        | 0.8107        | 0.6533        |
| LightGBM       | None             | 0.7077        | 0.6937        | 0.0173        | 0.7140        | 0.2542        | 0.5520        | 0.3481        | 0.9114        | 0.7400        | 0.8168        | 0.6455        |
| Random Forest  | None             | 0.6993        | 0.6964        | 0.0082        | 0.7979        | 0.3097        | 0.3756        | 0.3395        | 0.8962        | 0.8656        | 0.8807        | 0.5698        |
| Random Forest  | SMOTE-ENN        | 0.6866        | 0.6911        | 0.0113        | 0.7372        | 0.2506        | 0.4525        | 0.3226        | 0.8991        | 0.7829        | 0.8370        | 0.5930        |
| <b>XGBoost</b> | <b>SMOTE-ENN</b> | <b>0.6801</b> | <b>0.6857</b> | <b>0.0148</b> | <b>0.5551</b> | <b>0.1983</b> | <b>0.7285</b> | <b>0.3117</b> | <b>0.9237</b> | <b>0.5272</b> | <b>0.6713</b> | <b>0.6995</b> |
| LightGBM       | SMOTE-ENN        | 0.6757        | 0.6905        | 0.0167        | 0.7785        | 0.2616        | 0.3303        | 0.2920        | 0.8878        | 0.8504        | 0.8687        | 0.5376        |

*Note.* F = Failure class; S = Success class; AUC = test-set ROC-AUC; CV-AUC = 5-fold CV mean;  $\pm$ Std = CV standard deviation; Acc. = Accuracy; Prec. = Precision; Score = composite  $0.6 \times \text{AUC} + 0.4 \times \text{Recall (F)}$ . Bold row indicates recommended model.

## Model Performance Visualizations

**Feature set selection.** Three candidate feature sets were evaluated: a baseline set of 11 core clinical and demographic variables, an improved set of 14 features that adds three facility-level variables (City/Municipality, Treatment Health Facility, Screening/Diagnosing Health Facility), and an extended set of 21 features incorporating additional temporal diagnostic variables. Cross-validation performance consistently favored the improved set, which encodes geographic and facility-level care quality signals with minimal added complexity. The extended features did not yield consistent AUC gains and introduced variables with higher missingness rates. The improved 14-feature set was therefore used for all final model training and evaluation.


Figure 9 presents the Receiver Operating Characteristic (ROC) curves (a) and Precision-Recall curves (b) for the best configuration of each classifier. Panel (a) summarizes each model's ability to discriminate between treatment success and failure across decision thresholds, with the dashed diagonal representing random chance. Panel (b) provides a more informative view under class imbalance (approximately 86% success rate), where the horizontal dashed line marks the no-skill baseline at the failure class prevalence. Together, these two panels reveal both the overall discriminative capacity and the precision-recall trade-off that clinicians must navigate when setting alert thresholds.



**Figure 9**

(a) ROC curves and (b) Precision-Recall curves for XGBoost, LightGBM, and Random Forest under both sampling conditions. AUC and Average Precision (AP) values are shown in the legend. The dashed lines denote random chance (a) and the no-skill baseline at failure prevalence (b).


Figure 10 presents the confusion matrix for the best-performing model on the held-out test set, displaying the raw counts alongside row-normalized percentages. The matrix makes explicit the clinical trade-off between sensitivity (correctly identified failure cases) and specificity (correctly identified success cases), which is the primary decision axis for deploying the model in a CDSS alert system.



**Figure 10**

*Confusion matrix for the best-performing model on the test set. Raw counts are shown with row-normalized percentages in parentheses.*

Figure 11 presents the cross-validation stability of all six model configurations (c) and the top 15 feature importances from the best-performing classifier (d). Panel (c) shows the 5-fold CV ROC-AUC mean with standard deviation bars for each model–sampler combination, color-coded by sampling strategy; configurations with narrower error bars indicate more stable generalization. Panel (d) ranks the most predictive features by normalized importance, providing clinically interpretable signals for which patient attributes drive the risk prediction.



**Figure 11**

(c) 5-fold cross-validation ROC-AUC (mean  $\pm$  std) for all six model-sampler configurations, color-coded by sampling strategy. (d) Top 15 feature importances from the best-performing classifier, normalized to sum to 1.

Table 7 presents the full metric breakdown for all six configurations. Models were ranked by a composite clinical score  $S = 0.6 \times \text{AUC} + 0.4 \times \text{Recall}(\text{Failure})$ , which weights discriminative ability alongside sensitivity to treatment failure — the outcome that CDSS alerts must prioritize.

**XGBoost with SMOTE-ENN** was selected as the recommended model ( $S = 0.699$ ). It achieved the highest failure-class recall of any configuration (72.9%), meaning it correctly flagged nearly three-quarters of patients who ultimately failed treatment. Although its overall accuracy was lower (55.5%) due to SMOTE-ENN’s aggressive resampling of the training set, and its test-set AUC was moderate (0.680), this trade-off is clinically justified: in a TB CDSS, a missed failure (false negative) results in undetected treatment breakdown and potential disease progression or drug resistance, whereas a false alert triggers additional clinical review — a lower-cost error. The model’s 5-fold CV AUC of  $0.686 \pm 0.015$  confirmed stable generalization, and holdout validation yielded a recall of 68.2%, confirming the recall advantage persists on unseen data.

Among models without resampling, **XGBoost (no sampling)** achieved the highest test-set AUC (0.712) and CV-AUC ( $0.702 \pm 0.009$ ) across all six configurations, with a competitive failure recall of 56.6%. This configuration is preferred when a continuous, well-calibrated risk score is needed — for instance, to rank patients by risk level rather than apply a binary alert threshold. **LightGBM (no sampling)** performed comparably (AUC = 0.708, recall = 55.2%), providing a useful corroborating model. The three



configurations not selected — Random Forest (both sampling strategies) and LightGBM with SMOTE-ENN — showed lower composite scores, driven by weaker failure recall relative to their AUC costs.

## Temporal Models Performance

**Table 8**

*Version 1 (v1) Test Set Performance Metrics of Temporal Models for TB Treatment Outcome Prediction*

| Metric      | Bi-LSTM Baseline | Bi-LSTM Augmented | XGBoost Baseline | XGBoost SMOTE-ENN | LightGBM SMOTE-ENN | Random Forest SMOTE-ENN |
|-------------|------------------|-------------------|------------------|-------------------|--------------------|-------------------------|
| Accuracy    | 0.9048           | 0.9048            | 0.9048           | 0.9048            | 0.9048             | 0.7143                  |
| Precision   | 0.9048           | 0.9048            | 0.9048           | 0.9048            | 0.9048             | 0.8824                  |
| Recall      | 1.0000           | 1.0000            | 1.0000           | 1.0000            | 1.0000             | 0.7895                  |
| F1 Score    | 0.9500           | 0.9500            | 0.9500           | 0.9500            | 0.9500             | 0.8333                  |
| Specificity | 0.0000           | 0.0000            | 0.0000           | 0.0000            | 0.0000             | 0.0000                  |
| ROC-AUC     | 0.4737           | 0.3684            | 0.5789           | 0.1053            | 0.2105             | 0.5263                  |
| PR-AUC      | 0.8847           | 0.9099            | 0.9501           | 0.8356            | 0.8500             | 0.9419                  |
| Threshold   | 0.15             | 0.86              | 0.10             | 0.34              | 0.43               | 0.70                    |

Table 8 reports the v1 temporal-model results on the held-out test split. Although several models achieved high accuracy (0.9048) and F1 score (0.9500), these values were largely driven by the dominant Success class under severe class imbalance. The v1 specificity of 0.0000 indicates that none of the actual Failure cases were correctly identified on the test split, motivating the corrected v2 temporal modeling iteration.

**Table 9**

*Version 2 (v2) Test Set Performance Metrics of Temporal Models (M12)*

| Model                      | Accuracy | Bal. Acc. | Precision | Recall | F1     | Specificity | ROC-AUC | PR-AUC | Threshold |
|----------------------------|----------|-----------|-----------|--------|--------|-------------|---------|--------|-----------|
| Random Forest (+SMOTE-ENN) | 0.9318   | 0.8904    | 0.9730    | 0.9474 | 0.9600 | 0.8333      | 0.9693  | 0.9950 | 0.67      |
| LightGBM (+SMOTE-ENN)      | 0.9318   | 0.8904    | 0.9730    | 0.9474 | 0.9600 | 0.8333      | 0.8991  | 0.9779 | 0.78      |
| XGBoost (+SMOTE-ENN)       | 0.9318   | 0.8904    | 0.9730    | 0.9474 | 0.9600 | 0.8333      | 0.9298  | 0.9863 | 0.66      |
| XGBoost (Baseline)         | 0.9091   | 0.8772    | 0.9722    | 0.9211 | 0.9459 | 0.8333      | 0.9430  | 0.9895 | 0.81      |
| Hybrid Bi-LSTM (Baseline)  | 0.8409   | 0.9079    | 1.0000    | 0.8158 | 0.8986 | 1.0000      | 0.9868  | 0.9981 | 0.55      |
| Hybrid Bi-LSTM (Augmented) | 0.7500   | 0.8553    | 1.0000    | 0.7105 | 0.8308 | 1.0000      | 0.9474  | 0.9910 | 0.89      |

The v2 results (Table 9) reflect the corrected temporal modeling workflow after addressing data leakage risk, outcome-label handling, and improved temporal feature preparation. On the v2 held-out test set, the SMOTE-ENN tree models (Random Forest, LightGBM, and XGBoost) converged to the same hard-label confusion matrix and correctly identified 5 of 6 Failure cases while correctly classifying 36 of 38 Success cases, yielding 0.9318 accuracy and 0.8333 specificity. Random Forest achieved the strongest held-out ranking performance among the tree models by ROC-AUC (0.9693) and PR-AUC (0.9950). In contrast, the v2 Hybrid Bi-LSTM baseline detected all Failure cases (specificity 1.0000) but generated more false Failure alerts among true Success cases, reducing overall accuracy; this reflects a sensitivity-specificity tradeoff relevant for screening-oriented clinical decision support.

**Table 10**

*Version 2 (v2) Confusion Matrix Counts of Temporal Models*

| Model                      | TP | FN | FP | TN |
|----------------------------|----|----|----|----|
| Random Forest (+SMOTE-ENN) | 5  | 1  | 2  | 36 |
| LightGBM (+SMOTE-ENN)      | 5  | 1  | 2  | 36 |
| XGBoost (+SMOTE-ENN)       | 5  | 1  | 2  | 36 |
| XGBoost (Baseline)         | 5  | 1  | 3  | 35 |
| Hybrid Bi-LSTM (Baseline)  | 6  | 0  | 7  | 31 |
| Hybrid Bi-LSTM (Augmented) | 6  | 0  | 11 | 27 |


Note. Failure is treated as the positive class. TP = actual Failure predicted as Failure; FN = actual Failure predicted as Success; FP = actual Success predicted as Failure; TN = actual Success predicted as Success.



**Table 11***Methodological Differences Between Temporal Modeling Iterations (v1 vs v2)*

| Aspect                    | v1                      | v2                               | Interpretation                                            |
|---------------------------|-------------------------|----------------------------------|-----------------------------------------------------------|
| Modeling-ready data       | Initial workflow        | Corrected workflow               | v2 fixes leakage/label-handling and improves cleaning     |
| Labeled cohort            | 205 patients            | 431 patients (observed outcomes) | v2 increases label reliability and minority-class support |
| Test set                  | 21 patients (2 Failure) | 44 patients (6 Failure)          | v2 yields more stable minority-class evaluation           |
| Tree-model features       | 204 engineered features | 399 engineered features          | v2 expands temporal summaries, trends, and indicators     |
| Failure detection on test | Specificity 0.0000      | Specificity 0.8333–1.0000        | v2 meaningfully improves Failure identification           |

**Table 12***Model-Level Performance Comparison Between Temporal Model v1 and v2 (Test Sets)*


| Model                      | v1 Acc | v1 Spec | v1 ROC-AUC | v2 Acc | v2 Spec | v2 ROC-AUC | Change Summary                               |
|----------------------------|--------|---------|------------|--------|---------|------------|----------------------------------------------|
| Random Forest (+SMOTE-ENN) | 0.7143 | 0.0000  | 0.5263     | 0.9318 | 0.8333  | 0.9693     | Stronger Failure detection and ranking       |
| LightGBM (+SMOTE-ENN)      | 0.9048 | 0.0000  | 0.2105     | 0.9318 | 0.8333  | 0.8991     | Improved minority-class separation           |
| XGBoost (+SMOTE-ENN)       | 0.9048 | 0.0000  | 0.1053     | 0.9318 | 0.8333  | 0.9298     | Large ROC-AUC and specificity gain           |
| XGBoost (Baseline)         | 0.9048 | 0.0000  | 0.5789     | 0.9091 | 0.8333  | 0.9430     | Better calibrated Failure detection          |
| Hybrid Bi-LSTM (Baseline)  | 0.9048 | 0.0000  | 0.4737     | 0.8409 | 1.0000  | 0.9868     | Catches all Failures; more false alerts      |
| Hybrid Bi-LSTM (Augmented) | 0.9048 | 0.0000  | 0.3684     | 0.7500 | 1.0000  | 0.9474     | Higher sensitivity; reduced overall accuracy |

**Figure 12***Version 2 (v2) Hybrid Bi-LSTM Per-Month Performance (M0–M12)*

**Figure 13***Version 2 (v2) XGBoost Per-Month Performance (M0–M12)***Figure 14***Version 2 (v2) LightGBM Per-Month Performance (M0–M12)*

**Figure 15**

*Version 2 (v2) Random Forest Per-Month Performance (M0–M12)*




Figures 12–15 summarize v2 temporal performance as treatment monitoring data accumulates from M0 to M12 across the Hybrid Bi-LSTM, XGBoost, LightGBM, and Random Forest temporal models. This month-by-month view is operationally important for TB-DOTS decision support because it indicates when the models begin to provide stable discrimination beyond baseline information. The plots complement the single-point M12 metrics in Tables 9 and 10 by showing that temporal prediction quality evolves over the course of therapy rather than being a single end-of-treatment estimate.


### Correlation Analysis of Clinical and Temporal Features

**Figure 16**

*Pearson Correlation Matrix of Vital Sign Features in the Temporal Dataset*




Pairwise correlations among vital signs within the temporal monitoring dataset confirm that several clinical measurements are moderately associated, supporting the use of MICE. The strongest observed correlation was between age and heart rate ( $r = -0.32$ ), followed by oxygen saturation and heart rate ( $r = -0.30$ ), and systolic with diastolic blood pressure ( $r = 0.24$ ). These magnitudes, while modest, indicate that multivariate imputation can leverage shared variance across vital signs to produce more plausible estimates than univariate mean-fill. Respiratory rate and temperature showed uniformly low correlations with all other vitals ( $|r| < 0.10$ ), suggesting that their missingness is best handled by the missing indicator

**Figure 17***Categorical Feature Raw Cardinality — Justification for Stage 4 Harmonization*

with fill rather than MICE.

The cardinality audit of raw categorical features (Figure 17) confirms that data entry inconsistency was widespread across the temporal dataset. Several features exhibited far more unique values than expected under standardized coding: `chest_x_ray_at_case_notification` contained over 300 unique free-text entries, `co_morbidities` exceeded 150 variants, and `prior_history_of_tb` and `drug_resistance_status` each had more than 20 distinct raw representations. Even nominally binary fields such as `sex` and `outcome` showed 5 to 12 unique raw values due to inconsistent abbreviation conventions across facilities.

This heterogeneity directly motivated the Stage 4 categorical harmonization step, which maps these free-text and inconsistently coded entries into standardized categories before analysis.

**Figure 18***Cramér's V Correlation Matrix of Raw Categorical Variables in the Temporal Dataset*


Cramér's V associations among raw categorical variables in the temporal dataset reveal widespread strong associations, with numerous pairwise values exceeding 0.80. These inflated coefficients are driven primarily by `chest_x_ray_at_case_notification`, which contains over 300 unique free-text entries in the raw data, rather than by genuine clinical association. Similarly, `co_morbidities` contributed artificially elevated associations due to inconsistent coding across facilities. This finding directly motivated the Stage 4 categorical harmonization step, which standardizes free-text entries into clinically meaningful group-




**Figure 19***Cramér's V Correlation Matrix of Imputed Categorical Variables*

ings and eliminates spurious associations arising from data entry inconsistencies rather than true patient characteristics.

The imputed Cramér's V matrix (Figure 19) confirms that the preprocessing pipeline reduced spurious categorical associations, with median off-diagonal V dropping from approximately 0.46 to 0.07. Three columns from the raw matrix — `prior_history_of_tb`, `dat_supported`, and `nationality` — are absent because they lacked sufficient variance or non-missing pairs in the cleaned data to compute reliable associations.

**Figure 20***Temporal Feature Correlation Matrix of Treatment Variables*

The temporal feature correlation matrix (Figure 20) reveals the relational structure among key treatment variables aggregated over the M0–M7 window. Adherence metrics—including percentage\_adherence, doses\_taken\_per\_month, and cumulative\_doses\_taken—form a tight, strongly inter-correlated cluster ( $r > 0.85$ ), confirming internal measurement consistency. Anthropometric features (weight, height) are orthogonal to all adherence measures ( $r \sim 0$ ), indicating that body-size variation is independent of treatment compliance in this cohort. Retention features (last\_month\_attended, total\_observed\_months) show only weak positive associations with adherence ( $r \approx 0.2$ ), suggesting that visit persistence and dose compli-

**Figure 21***Temporal Lag Correlation of Percentage Adherence Across Consecutive Treatment Months*

ance capture distinct dimensions of patient engagement. Diagnostic variables such as `smear_microscopy_result` exhibit near-zero correlations with adherence and retention, consistent with their role as outcome indicators rather than behavioral predictors. Xpert MTB/RIF does not appear in the matrix because all observed values in the temporal subset were identical (zero variance), producing undefined Pearson coefficients.

**Table 13**

*Lag-1 Temporal Correlation of Pct Adherence Across Consecutive Treatment Months (Justifying Time-Series Model Architecture)*

| Lag Pair | Pearson r | p-value | Significant? | N (pairs) |
|----------|-----------|---------|--------------|-----------|
| M0 → M1  | 0.9242    | 0.0000  | Yes          | 300       |
| M1 → M2  | 1.0       | 0.0000  | Yes          | 285       |
| M2 → M3  | 1.0       | 0.0000  | Yes          | 275       |
| M3 → M4  | 1.0       | 0.0000  | Yes          | 271       |
| M4 → M5  | 0.9999    | 0.0000  | Yes          | 254       |
| M5 → M6  | 1.0       | 0.0000  | Yes          | 197       |
| M6 → M7  | 0.6202    | 0.0000  | Yes          | 80        |
| M7 → M8  | 0.2066    | 0.6945  | No           | 6         |

The month-to-month stability of adherence behavior in the temporal dataset provides empirical support for the choice of temporal model architecture. Early treatment months exhibited near-perfect lag-1 correlations (M0–M1:  $r = 0.92$ ,  $p < 0.001$ ; M1–M2:  $r = 1.00$ ,  $p < 0.001$ ), indicating that adherence is highly autocorrelated in the initial treatment phase. Correlation remained above  $r = 0.99$  through M5–M6,

then declined at M6–M7 ( $r = 0.62$ ,  $p < 0.001$ ) before becoming non-significant at M7–M8 ( $r = 0.21$ ,  $p = 0.69$ ), where only six patient pairs remained available. This sustained autocorrelation confirms that adherence carries meaningful sequential structure, justifying the use of recurrent architectures (RNN, LSTM) that model temporal dependencies (Nayebi et al., 2022). The decline after M6 also aligns with the retention cascade, suggesting that dropout-compromised patients exhibit increasingly erratic adherence before loss to follow-up.

## Discussion

This study investigated the use of machine learning models to predict treatment outcomes—Success or Failure—of PTB patients under the DOH-CAR TB-DOTS program, using monthly patient monitoring data alongside baseline clinical features. Temporal modeling was implemented in two iterations. The v1 iteration served as an initial baseline, but its results demonstrate a common pitfall in imbalanced clinical prediction: high accuracy can be achieved by favoring the majority Success class while failing to detect true treatment failures. This is reflected by v1 specificity of 0.0000 on the held-out test split, indicating that Failure identification was not reliable under the initial workflow.

The corrected v2 iteration improved the methodological defensibility of temporal prediction by addressing leakage risk, outcome-label handling, and temporal feature preparation. On the v2 held-out test split, the SMOTE-ENN tree models achieved consistent performance and identified 5 of 6 Failure cases while correctly classifying 36 of 38 Success cases. Random Forest exhibited the strongest held-out ranking performance among the tree models (ROC-AUC 0.9693, PR-AUC 0.9950), while LightGBM reported the strongest cross-validated balanced accuracy in the saved artifacts. The v2 Hybrid Bi-LSTM baseline detected all Failure cases on the held-out test split (specificity 1.0000) but produced more false Failure alerts among true Success cases, lowering overall accuracy. This tradeoff can be acceptable in a screening-oriented CDSS where missing high-risk patients is more costly than additional review workload, but it must be managed to avoid over-alerting.

The study contributes in several ways. By integrating temporal monitoring data with static baseline features, it demonstrates that patient-level longitudinal information improves predictive performance over static data alone. The application of SMOTE-ENN and advanced data augmentation techniques for Bi-LSTM

effectively addresses the extreme class imbalance inherent in TB treatment outcomes, enhancing the model's capacity to detect high-risk patients. Furthermore, the findings suggest that predictive models have practical value for clinical decision support systems, enabling clinicians to identify patients at risk of treatment failure early and prioritize interventions accordingly.

Based on these findings, it is recommended that healthcare facilities and policymakers consider integrating the v2 temporal modeling outputs into existing TB monitoring programs to improve early identification of patients at risk of adverse outcomes. For deployment-oriented performance, the v2 tree models (particularly Random Forest and LightGBM with SMOTE-ENN) provide a strong balance of Failure detection and overall correctness, while the Hybrid Bi-LSTM baseline may be useful when a more conservative screening stance is desired. Future research should focus on larger, multi-region datasets to enhance generalization, evaluate temporal calibration over time, and explore hybrid modeling approaches that combine deep temporal representations with tree-based learners. Additionally, enhancing interpretability of temporal dynamics (e.g., attention over treatment months) may improve clinician trust and adoption.

### **Safety Scope and Interpretation Limits**

The narrative explanation layer improves accessibility of model reasoning by translating contribution signals into clinician-readable text, but its outputs must be interpreted within clear safety boundaries. The explanation should be read as describing how factors increased or decreased the model-predicted risk, not as establishing causal mechanisms of disease progression. This distinction is important in high-stakes settings where generative outputs may otherwise be over-interpreted beyond their evidentiary basis (Vasey et al., 2022; World Health Organization, 2021).

Accordingly, the system is positioned as decision support: it summarizes model behavior while preserving clinician authority for diagnosis and treatment decisions. This framing is consistent with contemporary reporting and governance guidance for AI-assisted clinical prediction systems, where transparency, bounded use, and human oversight remain central requirements (Collins et al., 2024).

### **Acknowledgments**

The authors would like to express their sincere gratitude to Saint Louis University for providing the academic environment and resources that made this research possible. We extend our appreciation to the School of Accountancy, Management, Computing, and Information Studies (SAMCIS), specifically the Department of Computing and Information Studies, for their continuous support and guidance throughout the study. Our deepest thanks go to Dr. Randy Domantay, whose leadership and encouragement have been invaluable. We are especially grateful to our research adviser, Dr. Josephine S. Dela Cruz, for their expert guidance, insightful feedback, and unwavering support at every stage of the research. Lastly, we acknowledge the collaboration and contributions of our partner researchers from the SONAHBS Department of Pharmacy, whose expertise and assistance significantly enriched this study.

## References

- Alruhaymi, A., & Kim, C. (2021). Why can multiple imputations and how (mice) algorithm work? *Open Journal of Statistics*, 11(5), 759–777. <https://doi.org/10.4236/ojs.2021.115045>
- Amante, T. D., & Abdosh, T. (2015). Risk factors for unsuccessful tuberculosis treatment outcome (failure, default and death) in public health institutions, eastern ethiopia. *Pan African Medical Journal*, 20. <https://doi.org/10.11604/pamj.2015.20.247.3345>
- Aracri, F., Bianco, M. G., Quattrone, A., & Sarica, A. (2025). Bridging the gap: Missing data imputation methods and their effect on dementia classification performance. *Brain Sciences*, 15(6), 639. <https://doi.org/10.3390/brainsci15060639>
- Austin, P. C., White, I. R., Lee, D. S., & van Buuren, S. (2021). Missing data in clinical research: A tutorial on multiple imputation. *Canadian Journal of Cardiology*, 37(9), 1322–1331. <https://doi.org/10.1016/j.cjca.2020.11.010>
- Azeez, A., Ndege, J., & Mutambayi, R. (2018). Associated factors with unsuccessful tuberculosis treatment outcomes among tuberculosis/hiv coinfectd patients with drug-resistant tuberculosis. [https://journals.lww.com/ijmy/fulltext/2018/07040/Associated\\_Factors\\_with\\_Unsuccessful\\_Tuberculosis.8.aspx](https://journals.lww.com/ijmy/fulltext/2018/07040/Associated_Factors_with_Unsuccessful_Tuberculosis.8.aspx)
- Baron, G. (2025). Many tb patients left untreated. <https://tribune.net.ph/2025/03/29/many-tb-patients-left-untreated>
- Batista, G. E. A. P. A., Prati, R. C., & Monard, M. C. (2004). A study of the behavior of several methods for balancing machine learning training data. *ACM SIGKDD Explorations Newsletter*, 6(1), 20–29. <https://doi.org/10.1145/1007730.1007735>
- Bendre, A. D., Peters, P. J., & Kumar, J. (2021). Tuberculosis: Past, present and future of the treatment and drug discovery research. *Current Research in Pharmacology and Drug Discovery*, 2, 100037. <https://doi.org/10.1016/j.crphar.2021.100037>
- Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32. <https://doi.org/10.1023/A:1010933404324>

- Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321–357. <https://doi.org/10.1613/jair.953>
- Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>
- Collins, G. S., Moons, K. G. M., Dhiman, P., Riley, R. D., Beam, A. L., Calster, B. V., Ghassemi, M., Liu, X., Reitsma, J. B., van Smeden, M., Boulesteix, A.-L., Camaradou, J., Celi, L. A., Denaxas, S., Denniston, A. K., Glocker, B., Golub, R. M., Harvey, H., Heinze, G., et al. (2024). TRIPOD+AI statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods. *BMJ*, 385, e078378. <https://doi.org/10.1136/bmj-2023-078378>
- Department of Health [DOH]. (2014). *National tb control program manual of procedures (5th ed.)* Department of Health. [https://www.philhealth.gov.ph/partners/providers/pdf/NTCP\\_MoP2014.pdf](https://www.philhealth.gov.ph/partners/providers/pdf/NTCP_MoP2014.pdf)
- Department of Health - Republic of the Philippines. (2022). Doh-car. doh-cordillera, usaid, recalibrate plan to detect and mitigate tuberculosis hot spots. <https://caro.doh.gov.ph/>
- Dong, Y., & Peng, C.-Y. J. (2013). Principled missing data methods for researchers. *SpringerPlus*, 2(1), 222. <https://doi.org/10.1186/2193-1801-2-222>
- Emmanuel, T., Maupong, T., Mpoeleng, D., Semong, T., Mphago, B., & Tabona, O. (2021). A survey on missing data in machine learning. *Journal of Big Data*, 8(1), 140. <https://doi.org/10.1186/s40537-021-00516-9>
- Florentino, J. L., Arao, R. M. L., Garfin, A. M. C., Gaviola, D. M. G., Tan, C. R., Yadav, R. P., Hiatt, T., Morishita, F., Siroka, A., Yamanaka, T., & Nishikiori, N. (2022). Expansion of social protection is necessary towards zero catastrophic costs due to tb: The first national tb patient cost survey in the philippines. *PLOS One*, 17(2), e0264689. <https://doi.org/10.1371/journal.pone.0264689>



- Groenwold, R. H. H., White, I. R., Donders, A. R. T., Carpenter, J. R., Altman, D. G., & Moons, K. G. M. (2012). Missing covariate data in clinical research: When and when not to use the missing-indicator method for analysis. *Canadian Medical Association Journal*, 184(11), 1265–1269. <https://doi.org/10.1503/cmaj.110977>
- Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The elements of statistical learning: Data mining, inference, and prediction* (2nd). Springer. <https://doi.org/10.1007/978-0-387-84858-7>
- Hinay, A. A., Mamalintaw, M. A., Damasin, M. L., Shamera, J., Adrian, B., Joy, C., Evan, L., Insular, L. C., Joy, A., Cadotdot, M. T., Elipio, M. R., Ashley, J., & Dayaganon, J. B. (2025). Sociodemographic and clinical predictors of tuberculosis and unsuccessful treatment outcomes in davao city, philippines: A retrospective cohort study. *International Journal of Environmental Research and Public Health*, 22(7), 1154–1154. <https://doi.org/10.3390/ijerph22071154>
- Junaid, K. P., Kiran, T., Gupta, M., Kishore, K., & Siwatch, S. (2025). How much missing data is too much to impute for longitudinal health indicators? A preliminary guideline for the choice of the extent of missing proportion to impute with multiple imputation by chained equations. *Population Health Metrics*, 23(1), 2. <https://doi.org/10.1186/s12963-025-00364-2>
- Kaur, R., & Sharma, A. (2021). An accurate integrated system to detect pulmonary and extra pulmonary tuberculosis using machine learning algorithms. *Inteligencia Artificial*, 24(68), 104–122. <https://doi.org/10.4114/intartif.vol24iss68pp104-122>
- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T.-Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems*, 30, 3146–3154.
- Khalid, T. N. F. T., Mohammad, W. M. Z. W., Samat, R. A., & Husain, N. R. N. (2022). Predictors of tuberculosis disease in smokers. *PeerJ*, 10, e13984. <https://doi.org/10.7717/peerj.13984>
- Kossakov, M., Kulichenko, D., Butusov, D., & Popov, A. (2024). Quantitative comparison of machine-learning clustering methods for tuberculosis data analysis. *Engineering Proceedings*, 20(1), 60. <https://doi.org/10.3390/engproc2024060020>

- Laine, S., & Aila, T. (2016). Temporal ensembling for semi-supervised learning. <https://arxiv.org/abs/1610.02242>
- Lamsal, D. K., Lewis, O. D., Smith, S., & Jha, N. (2009). Factors related to defaulters and treatment failure of tuberculosis patients in the DOTS program in the Sunsari district of Eastern Nepal. *SAARC Journal of Tuberculosis, Lung Diseases and HIV/AIDS*, 6(1), 25–30.
- Lu, J. Y., Zhu, J., Zhu, J., & Duong, T. Q. (2022). Long-short-term memory machine learning of longitudinal clinical data accurately predicts acute kidney injury onset in covid-19. *International Journal of Infectious Diseases*, 122, 802–810. <https://doi.org/10.1016/j.ijid.2022.07.034>
- Lundberg, S. M., & Lee, S.-I. (2017a). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765–4774.
- Lundberg, S. M., & Lee, S.-I. (2017b). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765–4774. <https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions>
- Nayebi, A., Tipirneni, S., Foreman, B., Ratcliff, J., Reddy, C. K., & Subbian, V. (2022). Recurrent neural network based time-series modeling for long-term prognosis following acute traumatic brain injury. <https://pmc.ncbi.nlm.nih.gov/articles/PMC8861707/>
- Ndambuki, J., Nzomo, J., Muregi, L., Mutuku, C., Makokha, F., Nthusi, J., Ambale, C., Lynen, L., & Decroo, T. (2020). Comparison of first-line tuberculosis treatment outcomes between previously treated and new patients: A retrospective study in machakos subcountry, kenya. *International Health*. <https://academic.oup.com/inthealth/article/13/3/272/5898744?login=false>
- Nishat, M. M., Faisal, F., Ratul, I. J., Al-Monsur, A., Ar-Rafi, A. M., Nasrullah, S. M., & Khan, M. R. H. (2022). A comprehensive investigation of the performances of different machine learning classifiers with SMOTE-ENN oversampling technique and hyperparameter optimization for imbalanced heart failure dataset. *Scientific Programming*, 2022, 3649406. <https://doi.org/10.1155/2022/3649406>

- Noubiap, J. J., Nansseu, J. R., Nyaga, U. F., Nkeck, J. R., Endomba, F. T., Kaze, A. D., Agbor, V. N., & Bigna, J. J. (2019). Global prevalence of diabetes in active tuberculosis: A systematic review and meta-analysis of data from 2.3 million patients with tuberculosis. *The Lancet Global Health*, 7(4), e448–e460. [https://doi.org/10.1016/S2214-109X\(18\)30487-X](https://doi.org/10.1016/S2214-109X(18)30487-X)
- Peng, A., Kong, X.-H., Liu, S., Zhang, H., Xie, L., Ma, L., Zhang, Q., & Chen, Y. (2024). Explainable machine learning for early predicting treatment failure risk among patients with tb-diabetes comorbidity. *Scientific Reports*, 14(1), 7567. <https://doi.org/10.1038/s41598-024-57446-8>
- Rajpurkar, P., Chen, E., Banerjee, O., & Topol, E. J. (2022). Ai in health and medicine. *Nature Medicine*, 28(1), 31–38. <https://doi.org/10.1038/s41591-021-01614-0>
- Salmi, M., Atif, D., Oliva, D., Abraham, A., & Ventura, S. (2024). Handling imbalanced medical datasets: Review of a decade of research. *Artificial Intelligence Review*, 57(10), 273. <https://doi.org/10.1007/s10462-024-10884-2>
- Sellergren, A., Kazemzadeh, S., Jaroensri, T., Kiraly, A., Traverse, M., Kohlberger, T., Xu, S., Jamil, F., Hughes, C., Lau, C., et al. (2025). Medgemma technical report. <https://arxiv.org/abs/2507.05201>
- Shickel, B., Tighe, P. J., Bihorac, A., & Rashidi, P. (2018). Deep EHR: A survey of recent advances in deep learning techniques for electronic health record (EHR) analysis. *IEEE Journal of Biomedical and Health Informatics*, 22(5), 1589–1604. <https://doi.org/10.1109/JBHI.2017.2767063>
- Siddiqui, A. N., Khayyam, K. U., & Sharma, M. (2016). Effect of diabetes mellitus on tuberculosis treatment outcome and adverse reactions in patients receiving directly observed treatment strategy in india. *BioMed Research International*, 1–11. <https://doi.org/10.1155/2016/7273935>
- Sisk, R., Sperrin, M., Peek, N., van Smeden, M., & Martin, G. P. (2023). Imputation and missing indicators for handling missing data in the development and deployment of clinical prediction models: A simulation study. *Statistical Methods in Medical Research*, 32(8), 1461–1477. <https://doi.org/10.1177/09622802231165001>

- Smith, I. D. M., Jones, G. T., Simpson, C. R., Smith, B. H., & Lee, R. J. (2021). Ensemble machine learning models for predicting chronic disease outcomes in electronic health records. *Medical Decision Making*, 41(2), 130–141. <https://doi.org/10.1177/0272989X20982548>
- Stop TB Partnership. (2021). Tb and covid: One year on - media brief. [https://stoptb.org/webadmin/cms/docs/20210316\\_TB%20and%20COVID\\_One%20Year%20on\\_Media%20Brief\\_FINAL.pdf](https://stoptb.org/webadmin/cms/docs/20210316_TB%20and%20COVID_One%20Year%20on_Media%20Brief_FINAL.pdf)
- The Borgen Project. (2025). Tb in the philippines: A persistent public health crisis. <https://borgenproject.org/tb-in-the-philippines/>
- Ting, K. M. (2011). Confusion matrix. In C. Sammut & G. I. Webb (Eds.), *Encyclopedia of machine learning* (pp. 209–209). Springer. [https://doi.org/10.1007/978-0-387-30164-8\\_157](https://doi.org/10.1007/978-0-387-30164-8_157)
- Tobin, E. H., & Tristram, D. (2024). Tuberculosis overview. In *Statpearls*. StatPearls Publishing. <http://www.ncbi.nlm.nih.gov/books/NBK441916/>
- Tok, P. S. K., Liew, S. M., Wong, L. P., Razali, A., Loganathan, T., Chinna, K., Ismail, N., & Kadir, N. A. (2020). Determinants of unsuccessful treatment outcomes and mortality among tuberculosis patients in malaysia. *PLOS One*. <https://doi.org/10.1371/journal.pone.0231986>
- Vasey, B., Nagendran, M., Campbell, B., Clifton, D. A., Collins, G. S., Denaxas, S., Denniston, A. K., Faes, L., Geerts, B., Ibrahim, M., Liu, X., Mateen, B. A., McCradden, M. D., Morgan, L., Ordish, J., Rogers, C., Saria, S., Ting, D. S. W., Watkinson, P., ... the DECIDE-AI Expert Group. (2022). Reporting guideline for the early-stage clinical evaluation of decision support systems driven by artificial intelligence: DECIDE-AI. *Nature Medicine*, 28, 924–933. <https://doi.org/10.1038/s41591-022-01772-9>
- von Hippel, P. T. (2018). How many imputations do you need? A two-stage calculation using a quadratic rule. *Sociological Methods & Research*, 49(3), 699–718. <https://doi.org/10.1177/0049124117747303>
- World Health Organization. (2021). *Ethics and governance of artificial intelligence for health: WHO guidance*. <https://www.who.int/publications/i/item/9789240029200>

- World Health Organization. (2022). *Global tuberculosis report 2021*. <https://www.who.int/publications/i/item/9789240037021>
- World Health Organization. (2023). *Global tuberculosis report 2023* (tech. rep.). World Health Organization. Geneva. <https://www.who.int/publications/i/item/9789240083851>
- Xu, G., Hu, X., Lian, Y., & Li, X. (2023). Diabetes mellitus affects the treatment outcomes of drug-resistant tuberculosis. <https://doi.org/10.1186/s12879-023-08765-0>
- Zeng, X., & Zhu, K. (2024). Enhancing the interpretability of shap values using large language models. <https://arxiv.org/abs/2409.00079>

## **Appendix A: iCreate Supplementary Material**

The following pages reproduce the original iCreate 2026 supplementary material in its source format for archival completeness.

# Treatment Success and Failure Prediction using Machine Learning Models for Pulmonary Tuberculosis Patients under the DOH-CAR TB-DOTS Program

## Author Names

Aaron Kyle M. Aguilar<sup>1</sup>, Mark Lestat C. Agustin<sup>2</sup>, Keanu Sonn M. Fortaleza<sup>3</sup>, Benny Gil A. Lactaotao<sup>4</sup>,  
Marven Joffre G. Luis<sup>5</sup>, Georcelle Leisley U. Nuarin<sup>6</sup>, Sean Drei A. Octavo<sup>7</sup>, Gebreyi Isaac P. Rabang<sup>8</sup>

## Affiliations

<sup>1 2 3 4 5 6 7 8</sup> Department of Computing and Information Studies, Saint Louis University, Baguio City, Philippines

## Corresponding Author:

Josephine S. Dela Cruz, DIT – jdelacruz@slu.edu.ph

## Abstract

This study proposes a multi-layer analytical framework using machine learning to predict pulmonary tuberculosis (PTB) treatment success or failure in the Cordillera Administrative Region (CAR), Philippines. Utilizing demographic, clinical, and diagnostic data from 2015 to 2025, supervised models (Logistic Regression, Support Vector Machines, and Extreme Gradient Boosting) and temporal models (Recurrent Neural Networks and Long Short-Term Memory) are trained to identify risk factors. Shapley Additive Explanations (SHAP) ensure interpretability. The models will be integrated into a clinical decision support tool for early intervention. Aligning with the Sustainable Development Goals for good health and well-being, this predictive tool aims to enhance TB control efforts, optimize resource allocation, and improve patient management.

## Keywords

*Clinical decision support, machine learning, predictive modeling, pulmonary tuberculosis, treatment outcomes*

## 1. Introduction

Pulmonary Tuberculosis (PTB) is a chronic infectious disease that continues to challenge public health due to its gradual progression and risk of treatment failure (Tobin & Tristram, 2024). In the Philippines, TB remains a major public health concern, with the country being one of the top contributors to the global TB burden (World Health Organization, 2022). In the Cordillera Administrative Region (CAR), the Department of Health (DOH) reported 437 confirmed TB cases in Baguio City alone in 2022, alongside an estimated 4,405 missing or unreported cases (Department of Health Republic of the Philippines, n.d.).

While multiple studies have examined risk factors such as sociodemographic characteristics, comorbidities, and medical history, findings remain fragmented as traditional statistical analyses often assume linear relationships and examine only a few factors simultaneously (Amante &

Abdosh, 2015). To address this gap, this study adopts an integrative machine learning framework to combine multiple dimensions of data in analyzing interrelated variables.

The specific research objectives of this study are as follows:

1. **Data Pre-processing & Feature Engineering:** To clean, pre-process, and digitalize PTB patient records, creating engineered features that ensure the dataset is optimized for statistical analysis and machine learning applications.
2. **Statistical Analysis & Profiling:** To perform descriptive and inferential statistics to profile the demographic, clinical, and treatment characteristics of the study population, identifying baseline variables that significantly influence treatment outcomes.
3. **AI Modeling & Comparative Analysis:** To develop, validate, and evaluate both non-temporal and temporal machine learning models for predicting PTB treatment outcomes, comparing the predictive performance of static baseline data against dynamic, time-dependent patient patterns throughout the course of therapy.
4. **Data Visualization & Clinical Interpretation:** To visualize patient risk levels, statistical trends, and predictive outputs from the AI models, providing clear insights to support clinical interpretation and the early identification of individuals at risk of poor outcomes.

## 2. Methodology

This study utilizes a retrospective dataset of PTB patients from the CAR covering the period from 2015 to 2025, consolidated by the DOH-CAR. Data undergoes a standardized preprocessing pipeline, including schema harmonization, error correction, and missing value imputation via Multivariate Imputation by Chained Equations, or MICE (Alruhaymi & Kim, 2021).

To predict treatment outcomes defined by the National Tuberculosis Program criteria, supervised models including Logistic Regression, Support Vector Machines (SVM), and Extreme Gradient Boosting (XGBoost) are utilized. Furthermore, temporal models using Recurrent Neural Networks (RNN) and Long Short-Term Memory (LSTM) are employed to capture longitudinal treatment dynamics (Lu et al., 2022; Nayebi et al., 2022). Model performance is assessed using standard classification metrics, and Shapley Additive Explanations (SHAP) are applied to ensure clinical interpretability (Peng et al., 2024).

## 3. Results and Discussion

This section presents the results of machine learning analyses aimed at predicting treatment success and failure among pulmonary tuberculosis patients under the DOH-CAR TB-DOTS program. Descriptive statistics summarize patient demographics, clinical characteristics, and treatment outcomes, while predictive model performance is evaluated using test and validation datasets. Both static (Random Forest, Gradient Boosting, LightGBM, XGBoost, Logistic Regression) and temporal models are compared, and feature importance analyses provide insights into the key factors influencing treatment outcomes. The discussion interprets these findings in



the context of clinical relevance and public health priorities, highlighting implications for early intervention, resource allocation, and strengthened TB management.

### 3.1 Dataset Characteristics and Class Distribution

The final dataset exhibited significant class imbalance, with 86.2% of cases classified as treatment success and 13.8% as treatment failure. To ensure robust evaluation, the dataset was stratified into 70% training, 20% testing, and 10% validation subsets. Feature engineering expanded the dataset to 15 predictors, including four newly introduced geographic and facility-level variables and one derived clinical delay variable (Diagnosis\_to\_Treatment\_days). Resampling using SMOTE combined with Edited Nearest Neighbors (SMOTE+ENN) was applied exclusively to the training set, increasing the minority (failure) class by 433% and achieving a near-balanced 0.70:1 ratio. This approach mitigated class bias while preserving realistic evaluation conditions in the test and validation sets.

### 3.2 Validation Set Performance and Generalization

**Table 1: Test Set Model Performance**

| Model             | AUC (Area Under the Curve) | Recall (Failure) | F1 (Failure) |
|-------------------|----------------------------|------------------|--------------|
| Random Forest     | 0.7154                     | 57.98%           | 0.3539       |
| Gradient Boosting | 0.7121                     | 44.68%           | 0.3262       |
| LightGBM          | 0.7119                     | 48.40%           | 0.3467       |
| XGBoost           | 0.7095                     | 47.87%           | 0.3416       |

Among all evaluated models, Random Forest achieved the highest overall performance on the test set, with an AUC of 0.7154 and an F1-score of 0.3539 for the failure class. It also demonstrated balanced recall (57.98%) without excessive false positives. Gradient Boosting, LightGBM, and XGBoost achieved comparable AUC values (~0.71), though with slightly lower recall for treatment failure. Logistic Regression demonstrated the highest recall (77.13%); however, its precision was notably low (19.6%), resulting in over-alerting and reduced overall accuracy (~51%). These findings suggest that while Logistic Regression aggressively identifies high-risk patients, it lacks specificity, potentially burdening healthcare systems with excessive false alarms.

On the 10% held-out validation set, Random Forest maintained superior performance, achieving a validation AUC of 0.7380 and recall of 52.13% for treatment failure. Notably, all models demonstrated positive AUC gains from test to validation sets, indicating strong generalization and absence of overfitting. The consistency of Random Forest across both evaluation phases confirms its robustness and suitability for deployment within a clinical decision support framework.

### 3.3 Validation Set Performance and Generalization

**Table 2: Test Set Performance Metrics of Temporal Models for TB Treatment Outcome Prediction**

| Metric      | Bi-LSTM Baseline | Bi-LSTM Augmented | XGBoost Baseline | XGBoost SMOTE-ENN | LightGBM SMOTE-ENN | Random Forest SMOTE-ENN |
|-------------|------------------|-------------------|------------------|-------------------|--------------------|-------------------------|
| Accuracy    | 0.9048           | 0.9048            | 0.9048           | 0.9048            | 0.9048             | 0.7143                  |
| Precision   | 0.9048           | 0.9048            | 0.9048           | 0.9048            | 0.9048             | 0.8824                  |
| Recall      | 1.0000           | 1.0000            | 1.0000           | 1.0000            | 1.0000             | 0.7895                  |
| F1 Score    | 0.9500           | 0.9500            | 0.9500           | 0.9500            | 0.9500             | 0.8333                  |
| Specificity | 0.0000           | 0.0000            | 0.0000           | 0.0000            | 0.0000             | 0.0000                  |
| ROC-AUC     | 0.4737           | 0.3684            | 0.5789           | 0.1053            | 0.2105             | 0.5263                  |
| PR-AUC      | 0.8847           | 0.9099            | 0.9501           | 0.8356            | 0.8500             | 0.9419                  |
| Threshold   | 0.15             | 0.86              | 0.10             | 0.34              | 0.43               | 0.70                    |

Table 1 presents the performance metrics of all temporal models evaluated on the held-out test set. Both Bi-LSTM variants and XGBoost Baseline achieved the highest overall accuracy (0.9048) and F1 score (0.9500), reflecting their ability to correctly predict the dominant Success class. Precision values were similarly high (0.9048), indicating few false positives for the majority class. However, recall for the minority Failure class reached 1.0 for these models, while specificity remained 0.0, demonstrating that no actual Failure cases were correctly identified in the small test set. Tree-based models augmented with SMOTE-ENN, particularly Random Forest, showed a more balanced performance. Random Forest achieved slightly lower overall accuracy (0.7143) but higher recall for the Failure class (0.7895) and competitive PR-AUC (0.9419), indicating better identification of high-risk patients. ROC-AUC scores further illustrate differences in probability calibration, with XGBoost Baseline achieving the highest ROC-AUC (0.5789), while the Bi-LSTM Augmented variant performed lower (0.3684), likely due to overfitting on limited Failure samples. Overall, these

results suggest that while Bi-LSTM and XGBoost excel at predicting the majority outcome, Random Forest with SMOTE-ENN offers the most effective detection of treatment failure, making it more suitable for clinical decision support applications.

#### 4. Conclusions and Recommendations

This study investigated the use of machine learning models to predict treatment outcomes—Success or Failure—of pulmonary tuberculosis patients under the DOH-CAR TB-DOTS program, using monthly patient monitoring data alongside baseline clinical features. Both temporal models, specifically the Bi-LSTM Baseline and Augmented variants, and tree-based models, including XGBoost, LightGBM, and Random Forest with SMOTE-ENN, were evaluated on a stratified 70/20/10 patient-level split, incorporating progressive temporal samples from treatment months M0 through M12. The test set performance revealed that while Bi-LSTM and XGBoost Baseline models achieved high overall accuracy (0.9048) and F1 scores (0.9500), their ability to detect treatment failures was limited, as indicated by zero specificity for the minority Failure class. In contrast, Random Forest with SMOTE-ENN demonstrated more balanced performance, achieving a recall of 0.7895 for treatment failure and a high PR-AUC (0.9419), indicating greater reliability in identifying patients at risk of adverse outcomes. These results highlight that tree-based models augmented with SMOTE-ENN are particularly effective at handling severe class imbalance, while temporal models provide strong predictive accuracy for the majority outcome.

The study contributes in several ways. By integrating temporal monitoring data with static baseline features, it demonstrates that patient-level longitudinal information improves predictive performance over static data alone. The application of SMOTE-ENN and advanced data augmentation techniques for Bi-LSTM effectively addresses the extreme class imbalance inherent in TB treatment outcomes, enhancing the model's capacity to detect high-risk patients. Furthermore, the findings suggest that predictive models, particularly Random Forest with SMOTE-ENN, have practical value for clinical decision support systems, enabling clinicians to identify patients at risk of treatment failure early and prioritize interventions accordingly.

Based on these findings, it is recommended that healthcare facilities and policymakers consider integrating machine learning models like Random Forest with SMOTE-ENN into existing TB monitoring programs to improve treatment outcomes and optimize resource allocation. Future research should focus on larger, multi-region datasets to enhance generalization, explore hybrid modeling approaches combining deep learning and tree-based methods, and investigate additional features such as social determinants of health and geographic risk factors. Additionally, enhancing model interpretability, particularly of Bi-LSTM attention outputs, could improve clinician trust and adoption.

Overall, this study aligns with Sustainable Development Goal 3 – Good Health and Well-Being by providing a data-driven approach to improving TB treatment outcomes. By accurately identifying patients at risk of treatment failure, these predictive models have the potential to reduce TB-related morbidity and mortality, inform targeted public health interventions, and contribute to more efficient and effective TB care within the Cordillera Administrative Region.

### Acknowledgments

The authors would like to express their sincere gratitude to Saint Louis University for providing the academic environment and resources that made this research possible. We extend our appreciation to the School of Accountancy, Management, Computing, and Information Studies (SAMCIS), specifically the Department of Computing and Information Studies, for their continuous support and guidance throughout the study. Our deepest thanks go to Dr. Randy Domantay, whose leadership and encouragement have been invaluable. We are especially grateful to our research adviser, Dr. Josephine S. Dela Cruz, for their expert guidance, insightful feedback, and unwavering support at every stage of the research. Lastly, we acknowledge the collaboration and contributions of our partner researchers from the SONAHBS Department of Pharmacy, whose expertise and assistance significantly enriched this study.

### References

- Alruhaymi, A., & Kim, C. (2021). Why can multiple imputations and how (MICE) algorithm work? *Open Journal of Statistics*, 11(5), 759–777. <https://doi.org/10.4236/ojs.2021.115045>
- Amante, T. D., & Abdosh, T. (2015). Risk factors for unsuccessful tuberculosis treatment outcome (failure, default and death) in public health institutions, Eastern Ethiopia. *Pan African Medical Journal*, 20. <https://doi.org/10.11604/pamj.2015.20.247.3345>
- Department of Health Republic of the Philippines. (n.d.). DOH-CAR. DOH-Cordillera, USAID, recalibrate plan to detect and mitigate tuberculosis hot spots. <https://caro.doh.gov.ph/%F0%9D%90%83%F0%9D%90%8E%F0%9D%90%87-%F0%9D%90%82%F0%9D%90%A8%F0%9D%90%AB%F0%9D%90%9D%F0%9D%90%A2%F0%9D%90%A5%F0%9D%90%A5%F0%9D%90%9E%F0%9D%90%AB%F0%9D%90%9A-%F0%9D%90%94%F0%9D%90%92%F0%9D%90%80/>
- Lu, J. Y., Zhu, J., Zhu, J., & Duong, T. Q. (2022). Long-short-term memory machine learning of longitudinal clinical data accurately predicts acute kidney injury onset in COVID-19: a two-center study. *International Journal of Infectious Diseases*, 122, 802–810. <https://doi.org/10.1016/j.ijid.2022.07.034>
- Nayebi, A., Tipirneni, S., Foreman, B., Ratcliff, J., Reddy, C. K., & Subbian, V. (2022). Recurrent Neural Network based Time-Series Modeling for Long-term Prognosis Following Acute Traumatic Brain Injury. <https://pmc.ncbi.nlm.nih.gov/articles/PMC8861707/>

- Peng, A., Kong, X.-H., Liu, S., Zhang, H., Xie, L., Ma, L., Zhang, Q., & Chen, Y. (2024). Explainable machine learning for early predicting treatment failure risk among patients with TB-diabetes comorbidity. *Scientific Reports*, 14(1), 7567. <https://doi.org/10.1038/s41598-024-57446-8>
- Tobin, E. H., & Tristram, D. (2024). Tuberculosis overview. In StatPearls. StatPearls Publishing. <http://www.ncbi.nlm.nih.gov/books/NBK441916/>
- World Health Organization. (2022). Global tuberculosis report 2021: supplementary material. World Health Organization. <https://www.who.int/publications/i/item/9789240037021>